

Colab Link:

<https://colab.research.google.com/drive/14wfN6q9cSq-w6vTqYDpSOqfuQAYxn3pD?usp=sharing>

#Project Theme: US Company Bankruptcy Prediction

Team A08: Chih-Hsin (Olivia) Peng, Gaurangi Agrawal, Pei-Hsin (Bonny) Yang, Yu-Chun (Lila) Su

##1. Business Problem

Our analysis involves a thorough examination of numerous feature variables that contain comprehensive information about financial statement indices, from Current Assets to Total Operating Expenses, which are essential in providing an in-depth understanding of a company's financial health. In addition to this, we incorporate the temporal dimension by including a year column, which enables us to unveil trends and conditions over time. Our primary objective is to leverage this extensive dataset to predict a company's financial stability or its likelihood of facing bankruptcy.

##2. Data Source + Variable Description

The dataset we have chosen is from Kaggle, and the company name in the dataset is related to the American public company, which is listed on the New York Stock Exchange and NASDAQ.

Variable Name	Description
company-name	Name of the company
year	Operating years of the company
status_label	Tells us whether the company has declared bankruptcy or not. 1 - Alive/ Not Bankrupt : 0 - Failed/ Bankrupt
X1	Current assets - All the assets of a company that are expected to be sold or used as a result of standard business operations over the next year.
X2	Cost of goods sold - The total amount a company paid as a cost directly related to the sale of products.
X3	Depreciation and amortization - Depreciation refers to the loss of value of a tangible fixed asset over time (such as property, machinery, buildings, and plant). Amortization refers to the loss of value of intangible assets over time.
X4	EBITDA - Earnings before interest, taxes, depreciation, and amortization. It is a measure of a company's overall financial performance, serving as an alternative to net income.
X5	Inventory - The accounting of items and raw materials that a company either uses in production or sells.
X6	Net Income - The overall profitability of a company after all expenses and costs have been deducted from total revenue.
X7	Total Receivables - The balance of money due to a firm for goods or services delivered or used but not yet paid for by

Variable Name	Description
	customers.
X8	Market value - The price of an asset in a marketplace. In this dataset, it refers to the market capitalization since companies are publicly traded in the stock market.
X9	Net sales - The sum of a company's gross sales minus its returns, allowances, and discounts.
X10	Total assets - All the assets, or items of value, a business owns.
X11	Total Long-term debt - A company's loans and other liabilities that will not become due within one year of the balance sheet date.
X12	EBIT - Earnings before interest and taxes.
X13	Gross Profit - The profit a business makes after subtracting all the costs that are related to manufacturing and selling its products or services.
X14	Total Current Liabilities - The sum of accounts payable, accrued liabilities, and taxes such as Bonds payable at the end of the year, salaries, and commissions remaining.
X15	Retained Earnings - The amount of profit a company has left over after paying all its direct costs, indirect costs, income taxes, and its dividends to shareholders.
X16	Total Revenue - The amount of income that a business has made from all sales before subtracting expenses. It may include interest and dividends from investments.
X17	Total Liabilities - The combined debts and obligations that the company owes to outside parties.
X18	Total Operating Expenses - The expenses a business incurs through its normal business operations.

##3. Who cares about this problem & why?

Knowing who is concerned about a particular issue offers different perspectives. Firstly, for a company, understanding its financial health guides its future plans. If a company faces money troubles, it might find ways to earn more, spend less on making things, or seek outside help, like selling company shares or bonds. This understanding helps companies change how they work to stay financially strong.

At the same time, investors focus on how a company's finances affect their investments. They want to put money where they think it'll grow. By looking at a company's past finances, investors can guess if it might close down or not make money. Using this info, they decide where to keep investing and where to pull back. So, knowing a company's financial situation is crucial for both the company itself and the people investing in it. Predicting if a company will do well or not guides the actions of both businesses and investors.

##4. Data Preprocessing

```

%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

#import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv('/content/american_bankruptcy.csv')

# rename the column's name
df.rename(columns={'status_label' : 'is_bankruptcy', 'X1': 'Current
assets', 'X2': 'Cost of goods sold', 'X3': 'Depreciation and
amortization', 'X4': 'EBITDA', 'X5': 'Inventory', 'X6': 'Net Income',
'X7': 'Total Receivables', 'X8': 'Market Value', 'X9': 'Net Sales',
'X10': 'Total Assets', 'X11': 'Total Long-term Debt', 'X12': 'EBIT',
'X13': 'Gross Profit', 'X14': 'Total Current Liabilities', 'X15':
'Retained Earnings', 'X16': 'Total Revenue', 'X17': 'Total
Liabilities', 'X18': 'Total Operating Expenses'},inplace=True)

# set 'company_name' and 'year' as dual index
df.set_index(['company_name', 'year'],inplace=True)

# change 'is_bankruptcy' column into 0 = failed / 1 = alive
# set(df['status_label']) == {'alive', 'failed'}
mapping = {'alive':0, 'failed':1}
df['is_bankruptcy'] = df['is_bankruptcy'].replace(mapping)

```

##5. Feature Selection

```

# pre-processing the train/ test dataset, handling the missing value,
and standardize the variables
X = df.drop('is_bankruptcy',axis=1)
y = df['is_bankruptcy']

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =
train_test_split(X,y,test_size=0.25,random_state=42)

# dealing with null value before using ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler

numeric_pipeline =
Pipeline([('imputer',SimpleImputer(strategy='mean')),
('scaler',StandardScaler())])

```

```

                                # after handling the missing value,
                                standardize the numerical variables

from sklearn.compose import ColumnTransformer
numerical_variable = ['Current assets', 'Cost of goods sold',
                      'Depreciation and amortization', 'EBITDA', 'Inventory', 'Net
Income',
                      'Total Receivables', 'Market Value', 'Net Sales', 'Total
Assets',
                      'Total Long-term Debt', 'EBIT', 'Gross Profit',
                      'Total Current Liabilities', 'Retained Earnings', 'Total
Revenue',
                      'Total Liabilities', 'Total Operating Expenses']

transformer =
ColumnTransformer([('numerical_variable', numeric_pipeline, numerical_va
riable)])
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
pipeline = Pipeline([('transformer', transformer),
                     ('LogisticRegression', LogisticRegression(random_state=42))])
pipelineX = df.drop('is_bankruptcy', axis=1)
y = df['is_bankruptcy']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=0.25, random_state=42)

# dealing with null value before using ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler

numeric_pipeline =
Pipeline([('imputer', SimpleImputer(strategy='mean')),
         ('scaler', StandardScaler())])
                                # after handling the missing value,
                                standardize the numerical variables

from sklearn.compose import ColumnTransformer
numerical_variable = ['Current assets', 'Cost of goods sold',
                      'Depreciation and amortization', 'EBITDA', 'Inventory', 'Net
Income',
                      'Total Receivables', 'Market Value', 'Net Sales', 'Total
Assets',
                      'Total Long-term Debt', 'EBIT', 'Gross Profit',
                      'Total Current Liabilities', 'Retained Earnings', 'Total
Revenue',
                      'Total Liabilities', 'Total Operating Expenses']

```

```

transformer =
ColumnTransformer([('numerical_variable', numeric_pipeline, numerical_va
riable)])
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
pipeline = Pipeline([('transformer', transformer),
('LogisticRegression', LogisticRegression(random_state=42))])
pipeline

Pipeline(steps=[('transformer',
ColumnTransformer(transformers=[('numerical_variable',
Pipeline(steps=[('imputer',
SimpleImputer()),
('scaler',
StandardScaler())])),
['Current assets',
'Cost of goods
sold',
'Depreciation and '
'amortization',
'EBITDA',
'Inventory',
'Net Income',
'Total
Receivables',
'Market Value',
'Net Sales',
'Total Assets',
'Total Long-term
Debt',
'EBIT', 'Gross
Profit',
'Total Current
Liabilities',
'Retained
Earnings',
'Total Revenue',
'Total
Liabilities',
'Total Operating '
'Expenses']))]),
('LogisticRegression',
LogisticRegression(random_state=42))])

```

###5.1 Select KBest

```
# step 1: use SelectKBest to pick the best features
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import f_classif # because I get
numerical input and a categorical target variable, so choose f_classif
to make the calculation
best_k_select = SelectKBest(score_func=f_classif)

# step 2: re-define the pipeline you want to put in GridSearchCV
best_k_select_pipeline = Pipeline([('transformer',transformer),
('select',best_k_select),
('LogisticRegression',LogisticRegression(random_state=42))])

# step 3: use gridsearch to find the best features
from sklearn.model_selection import GridSearchCV
param_grid = {'select__k': range(1, X_train.shape[1] + 1)}
grid_search = GridSearchCV(estimator=best_k_select_pipeline,
param_grid=param_grid,cv=5, scoring='balanced_accuracy')
grid_search.fit(X_train,y_train)

# step 4: after fitting, show the result
# 4-1
result = pd.DataFrame(grid_search.cv_results_)
#display(grid_search_all_result)
result = result.sort_values(by='mean_test_score',ascending=False)
display(result.filter(regex= '(^param_|mean_test_score)', axis=1))
# regular expression
# ^ : start from
# | : or

# 4-2
final_pick = grid_search.best_estimator_
print('The selected features are=\n',
final_pick["transformer"].get_feature_names_out()
[final_pick["select"].get_support()])
```

	param_select__k	mean_test_score
17	18	0.503538
16	17	0.502628
15	16	0.502628
14	15	0.501163
13	14	0.501163
12	13	0.501163
11	12	0.500528
9	10	0.500334
10	11	0.500187
6	7	0.499957
2	3	0.499955
7	8	0.499947
5	6	0.499947

1	2	0.499946
8	9	0.499938
0	1	0.499927
3	4	0.499909
4	5	0.499900

The selected features are=

```
[ 'numerical_variable__Current assets'
'numerical_variable__Cost of goods sold'
'numerical_variable__Depreciation and amortization'
'numerical_variable__EBITDA' 'numerical_variable__Inventory'
'numerical_variable__Net Income' 'numerical_variable__Total
Receivables'
'numerical_variable__Market Value' 'numerical_variable__Net Sales'
'numerical_variable__Total Assets'
'numerical_variable__Total Long-term Debt' 'numerical_variable__EBIT'
'numerical_variable__Gross Profit'
'numerical_variable__Total Current Liabilities'
'numerical_variable__Retained Earnings'
'numerical_variable__Total Revenue'
'numerical_variable__Total Liabilities'
'numerical_variable__Total Operating Expenses']
```

###5.2 Select From Model

```
# step 1: decide use SelectFromModel to select the best features, and
decide to use Lasso model
from sklearn.feature_selection import SelectFromModel
from sklearn.linear_model import Lasso
select_from_model = SelectFromModel(Lasso())

# step 2: re-define the pipeline, so you can put in GridSearchCV
select_from_model_pipeline = Pipeline([('transformer',transformer),
('select',select_from_model),
('LogisticRegression',LogisticRegression(random_state=42))])

# step 3: use gridsearch to find the best features
from sklearn.model_selection import GridSearchCV
param_grid = {'select__estimator__alpha': [10, 100,
1000], 'select__threshold': ['mean', 'median']}
grid_search =
GridSearchCV(estimator=select_from_model_pipeline,param_grid =
param_grid, cv=5, scoring='balanced_accuracy')
grid_search.fit(X_train, y_train)

# step 4: after fitting, show the result
# 4-1
result = pd.DataFrame(grid_search.cv_results_)
result = result.sort_values(by="mean_test_score", ascending=False)
```

```
display(result.filter(regex='(^param_|mean_test_score)', axis=1))
# 4-2
final_pick = grid_search.best_estimator_
print('The selected features are=\n',
final_pick["transformer"].get_feature_names_out()
[final_pick["select"].get_support()])
```

	param_select__estimator__alpha	param_select__threshold	mean_test_score
0	10	mean	0.503538
1	10	median	0.503538
2	100	mean	0.503538
3	100	median	0.503538
4	1000	mean	0.503538
5	1000	median	0.503538

```
The selected features are=
['numerical_variable__Current assets'
'numerical_variable__Cost of goods sold'
'numerical_variable__Depreciation and amortization'
'numerical_variable__EBITDA' 'numerical_variable__Inventory'
'numerical_variable__Net Income' 'numerical_variable__Total
Receivables'
'numerical_variable__Market Value' 'numerical_variable__Net Sales'
'numerical_variable__Total Assets'
'numerical_variable__Total Long-term Debt' 'numerical_variable__EBIT'
'numerical_variable__Gross Profit'
'numerical_variable__Total Current Liabilities'
'numerical_variable__Retained Earnings'
'numerical_variable__Total Revenue'
'numerical_variable__Total Liabilities'
'numerical_variable__Total Operating Expenses']
```

###5.3 Recursive Feature Elimination

```
# step 1: decide to use RFECV to select the best features
from sklearn.feature_selection import RFECV
rfecv = RFECV(LogisticRegression(random_state=42),
scoring='balanced_accuracy')

# step 2: re-define the pipeline, add 'select' in to the previous
pipeline
rfecv_pipeline = Pipeline([('transformer',transformer),
```



```

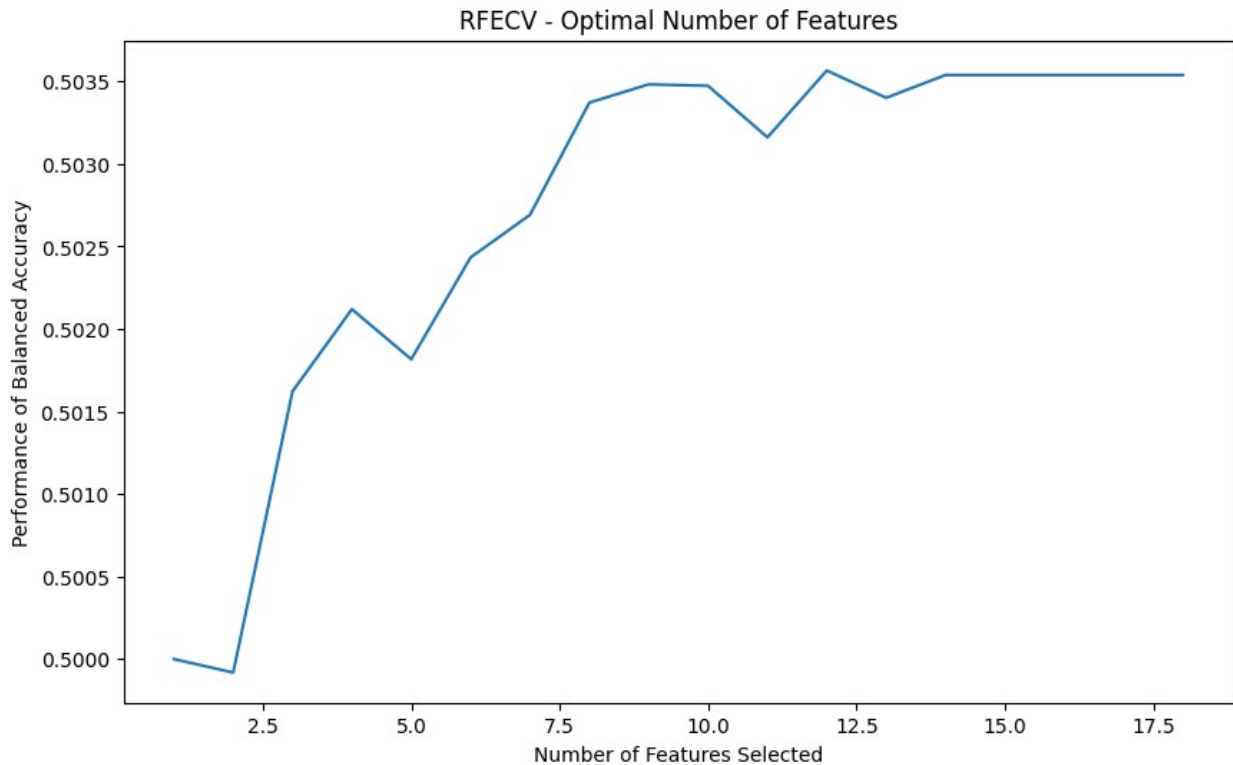
('select',rfecv),
('LogisticRegression',LogisticRegression(random_state=42))])
rfecv_pipeline.fit(X_train, y_train)

# step 3: show the result
print('Optimal number of features',rfecv.n_features_)
# .n_features_: attributes of rfecv, shows the number of selected
# features after the recursive feature elimination process
print('The selected features are=\n',
rfecv_pipeline['transformer'].get_feature_names_out()[rfecv.support_])

Optimal number of features 12
The selected features are=
['numerical_variable__Current assets' 'numerical_variable__EBITDA'
'numerical_variable__Inventory' 'numerical_variable__Net Income'
'numerical_variable__Total Receivables'
'numerical_variable__Market Value'
'numerical_variable__Total Long-term Debt' 'numerical_variable__EBIT'
'numerical_variable__Gross Profit'
'numerical_variable__Total Current Liabilities'
'numerical_variable__Retained Earnings'
'numerical_variable__Total Liabilities']

# step 4: plot
plt.figure(figsize=(10, 6))
plt.title("RFECV - Optimal Number of Features")
plt.xlabel("Number of Features Selected")
plt.ylabel("Performance of Balanced Accuracy")
mean_test_score = rfecv.cv_results_['mean_test_score']
plt.plot(range(1, len(mean_test_score) + 1), mean_test_score);

```



###5.4 Sequential Feature Selection: Forward and Backward

```
# default: forward
from mlxtend.feature_selection import SequentialFeatureSelector
sfs = SequentialFeatureSelector(LogisticRegression(random_state=42),
k_features="best", scoring='balanced_accuracy', forward=True)

sfs_pipeline = Pipeline([('transformer', transformer), ('select', sfs),
('LogisticRegression', LogisticRegression(random_state=42))])
sfs_pipeline.fit(X_train, y_train)

all_result = pd.DataFrame.from_dict(sfs.get_metric_dict()).T
display(all_result.sort_values(by='avg_score', ascending=False))
```

	feature_idx \
17	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, ...
16	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, ...
15	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, ...
14	(0, 1, 2, 3, 4, 5, 6, 7, 9, 10, 11, 12, 13, 16)
13	(0, 2, 3, 4, 5, 6, 7, 9, 10, 11, 12, 13, 16)
18	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, ...
9	(4, 5, 6, 7, 9, 10, 11, 12, 16)
11	(2, 3, 4, 5, 6, 7, 9, 10, 11, 12, 16)
10	(2, 4, 5, 6, 7, 9, 10, 11, 12, 16)
8	(4, 6, 7, 9, 10, 11, 12, 16)
12	(0, 2, 3, 4, 5, 6, 7, 9, 10, 11, 12, 16)

7	(4, 6, 7, 9, 10, 11, 16)
6	(6, 7, 9, 10, 11, 16)
5	(6, 9, 10, 11, 16)
4	(9, 10, 11, 16)
3	(9, 11, 16)
2	(11, 16)
1	(11,)

		cv_scores	avg_score	\
17	[0.5043231150613701, 0.5038300988665965, 0.503...	0.503...	0.503685	
16	[0.5043231150613701, 0.5038300988665965, 0.503...	0.503...	0.503685	
15	[0.5043231150613701, 0.5038300988665965, 0.503...	0.503...	0.503685	
14	[0.5043231150613701, 0.5038300988665965, 0.503...	0.503...	0.503685	
13	[0.5043231150613701, 0.5038300988665965, 0.503...	0.503...	0.503685	
18	[0.5035889196963023, 0.5038300988665965, 0.503...	0.503...	0.503538	
9	[0.5043231150613701, 0.5032311163304729, 0.502...	0.502...	0.503299	
11	[0.5043231150613701, 0.5032311163304729, 0.502...	0.502...	0.503299	
10	[0.5043231150613701, 0.5032311163304729, 0.502...	0.502...	0.503299	
8	[0.5043684624221537, 0.5033671584128236, 0.502...	0.502...	0.503224	
12	[0.5036796144178695, 0.5032311163304729, 0.503...	0.503...	0.503179	
7	[0.5037249617786531, 0.5034125057736073, 0.501...	0.501...	0.502994	
6	[0.5036796144178695, 0.5027681758767001, 0.501...	0.501...	0.502847	
5	[0.5029454190528017, 0.5027681758767001, 0.501...	0.501...	0.502304	
4	[0.5023926131308681, 0.5022145407013601, 0.501...	0.501...	0.501834	
3	[0.5005074585611499, 0.5004175930929893, 0.501...	0.501...	0.500684	
2	[0.4998186105568656, 0.5002815510106384, 0.501...	0.501...	0.500334	
1	[0.4998639579176492, 0.499773263196082, 0.5005...	0.5005...	0.500029	

std_dev	\	feature_names	ci_bound
17	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13,...	0.000553	0.00043
16	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13,...	0.000553	0.00043
15	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13,...	0.000553	0.00043
14	(0, 1, 2, 3, 4, 5, 6, 7, 9, 10, 11, 12, 13, 16)	0.000553	0.00043
13	(0, 2, 3, 4, 5, 6, 7, 9, 10, 11, 12, 13, 16)	0.000553	0.00043
18	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13,...	0.000333	0.000259
9	(4, 5, 6, 7, 9, 10, 11, 12, 16)	0.000886	0.000689
11	(2, 3, 4, 5, 6, 7, 9, 10, 11, 12, 16)	0.000886	0.000689
10	(2, 4, 5, 6, 7, 9, 10, 11, 12, 16)	0.000886	0.000689
8	(4, 6, 7, 9, 10, 11, 12, 16)	0.000861	

```

0.00067
12      (0, 2, 3, 4, 5, 6, 7, 9, 10, 11, 12, 16)  0.000642
0.0005
7      (4, 6, 7, 9, 10, 11, 16)  0.000864
0.000673
6      (6, 7, 9, 10, 11, 16)  0.000829
0.000645
5      (6, 9, 10, 11, 16)  0.000988
0.000768
4      (9, 10, 11, 16)  0.000548
0.000426
3      (9, 11, 16)  0.000739
0.000575
2      (11, 16)  0.000586
0.000456
1      (11,)  0.000352
0.000274

```

```

      std_err
17  0.000215
16  0.000215
15  0.000215
14  0.000215
13  0.000215
18  0.00013
9   0.000345
11  0.000345
10  0.000345
8   0.000335
12  0.00025
7   0.000336
6   0.000322
5   0.000384
4   0.000213
3   0.000287
2   0.000228
1   0.000137

```

back

```

from mlxtend.feature_selection import SequentialFeatureSelector
sbs = SequentialFeatureSelector(LogisticRegression(random_state=42),
k_features="best", scoring='balanced_accuracy', forward=False)

```

```

sbs_pipeline = Pipeline([('transformer', transformer), ('select', sbs),
('LogisticRegression', LogisticRegression(random_state=42))])
sbs_pipeline.fit(X_train, y_train)

```

```

all_result = pd.DataFrame.from_dict(sbs.get_metric_dict()).T
display(all_result.sort_values(by='avg_score', ascending=False))

```

	feature_idx \
9	(0, 3, 4, 6, 7, 10, 12, 13, 16)
16	(0, 1, 2, 3, 4, 6, 7, 8, 9, 10, 11, 12, 13, 15...
15	(0, 2, 3, 4, 6, 7, 8, 9, 10, 11, 12, 13, 15, 1...
14	(0, 2, 3, 4, 6, 7, 8, 9, 10, 11, 12, 13, 16, 17)
13	(0, 2, 3, 4, 6, 7, 8, 9, 10, 11, 12, 13, 16)
17	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13,...
10	(0, 3, 4, 6, 7, 9, 10, 12, 13, 16)
11	(0, 3, 4, 6, 7, 9, 10, 11, 12, 13, 16)
12	(0, 2, 3, 4, 6, 7, 9, 10, 11, 12, 13, 16)
18	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13,...
8	(0, 3, 4, 6, 7, 10, 13, 16)
7	(0, 3, 4, 7, 10, 13, 16)
6	(0, 3, 7, 10, 13, 16)
5	(0, 3, 7, 10, 16)
4	(0, 3, 7, 16)
3	(0, 7, 16)
1	(0,)
2	(0, 16)

	cv_scores	avg_score \
9	[0.5042777677005865, 0.5047918602889889, 0.502...	0.503785
16	[0.5049212683440871, 0.5040114883097309, 0.503...	0.503731
15	[0.5049212683440871, 0.5040114883097309, 0.503...	0.503731
14	[0.5049212683440871, 0.5040114883097309, 0.503...	0.503731
13	[0.5049212683440871, 0.5040114883097309, 0.503...	0.503731
17	[0.5043231150613701, 0.5038300988665965, 0.503...	0.503685
10	[0.5042777677005865, 0.5041475303920817, 0.502...	0.503665
11	[0.5042777677005865, 0.5040114883097309, 0.503...	0.503629
12	[0.5042777677005865, 0.5040114883097309, 0.503...	0.503611
18	[0.5035889196963023, 0.5038300988665965, 0.503...	0.503538
8	[0.5042777677005865, 0.5047011655674217, 0.502...	0.503491
7	[0.5036796144178695, 0.5041475303920817, 0.502...	0.503261
6	[0.5030361137743689, 0.5023052354229273, 0.502...	0.502764
5	[0.5030814611351524, 0.5023052354229273, 0.501...	0.502377
4	[0.5011056118438669, 0.5030856074021853, 0.501...	0.502102
3	[0.5005074585611499, 0.5024412775052781, 0.501...	0.501503
1	[0.5, 0.5, 0.5, 0.5, 0.5]	0.5
2	[0.5, 0.5, 0.5, 0.5, 0.4999546526392164]	0.499991

	feature_names	ci_bound
std_dev \		
9	(0, 3, 4, 6, 7, 10, 12, 13, 16)	0.001068
0.000831		
16	(0, 1, 2, 3, 4, 6, 7, 8, 9, 10, 11, 12, 13, 15...	0.001066
0.00083		
15	(0, 2, 3, 4, 6, 7, 8, 9, 10, 11, 12, 13, 15, 1...	0.001066
0.00083		
14	(0, 2, 3, 4, 6, 7, 8, 9, 10, 11, 12, 13, 16, 17)	0.001066
0.00083		

13	(0, 2, 3, 4, 6, 7, 8, 9, 10, 11, 12, 13, 16)	0.001066
0.00083		
17	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, ...)	0.000553
0.00043		
10	(0, 3, 4, 6, 7, 9, 10, 12, 13, 16)	0.000884
0.000688		
11	(0, 3, 4, 6, 7, 9, 10, 11, 12, 13, 16)	0.000838
0.000652		
12	(0, 2, 3, 4, 6, 7, 9, 10, 11, 12, 13, 16)	0.00086
0.000669		
18	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, ...)	0.000333
0.000259		
8	(0, 3, 4, 6, 7, 10, 13, 16)	0.001245
0.000968		
7	(0, 3, 4, 7, 10, 13, 16)	0.000963
0.000749		
6	(0, 3, 7, 10, 13, 16)	0.000682
0.00053		
5	(0, 3, 7, 10, 16)	0.000527
0.00041		
4	(0, 3, 7, 16)	0.001105
0.000859		
3	(0, 7, 16)	0.000863
0.000672		
1	(0,)	0.0
0.0		
2	(0, 16)	0.000023
0.000018		

	std_err
9	0.000415
16	0.000415
15	0.000415
14	0.000415
13	0.000415
17	0.000215
10	0.000344
11	0.000326
12	0.000335
18	0.00013
8	0.000484
7	0.000375
6	0.000265
5	0.000205
4	0.00043
3	0.000336
1	0.0
2	0.000009

In this project, although we considered models like KNN and Random Forest for training, our final choice for feature selection was logistic regression. Logistic regression was preferred for its simplicity, computational efficiency, interpretability, and its suitability for binary classification.

The process involved employing various methods for feature selection. SelectKBest recommended including all 18 features. SelectFromModel, using L1 regularization, also suggested including all features. Recursive feature elimination proposed retaining 12 features. Additionally, forward and backward sequential feature selection advised including 17 and 9 features, respectively. Surprisingly, most methods recommended retaining all features despite different approaches.

##6. Dataset Exploration

###Some Interesting Descriptive Analyses

We initiated our examination of the dataset, noting its emphasis on financial statements—spanning from Current Assets to Total Operating Expenses. This aligns well with our analysis objectives, leveraging the readily available and crucial financial data for a comprehensive understanding of a company's financial health. This is also the reason why the correlation between each variables is relatively high, as shown in correlation matrix. The inclusion of a year column allows us to uncover trends and conditions over time. If a company declared bankruptcy the previous year, it does not appear in the subsequent year.

Moving on to basic statistical analysis, we find the significantly larger number of alive companies compared to those declaring bankruptcy. This presents a notable challenge, which we will delve into in more detail in our subsequent analyses. We can also observe there is a substantial difference between alive companies and failed companies.

```
# show some basic information about the dataset
print('top 5 row of the dataset:')
display(df.head(5))
print()
print('info about the dataset:')
display(df.info())
print()
print('check if there are any NULL value in the dataset:')
display(df.isna().sum())
print()
print('basic statistics about the dataset (include = "number"):')
display(df.describe(include = 'number'))
print()
```

top 5 row of the dataset:

		is_bankruptcy	Current assets	Cost of goods
sold \ company_name	year			
C_1	1999	0	511.267	833.107
	2000	0	485.856	713.811

	2001	0	436.656	526.477
	2002	0	396.412	496.747
	2003	0	432.204	523.302
		Depreciation and amortization	EBITDA	Inventory \
company_name	year			
C_1	1999	18.373	89.031	336.018
	2000	18.577	64.367	320.590
	2001	22.496	27.207	286.588
	2002	27.172	30.745	259.954
	2003	26.680	47.491	247.245
		Net Income	Total Receivables	Market Value Net
Sales \	company_name	year		
C_1	1999	35.163	128.348	372.7519
1024.333	2000	18.531	115.187	377.1180
874.255	2001	-58.939	77.528	364.5928
638.721	2002	-12.410	66.322	143.3295
606.337	2003	3.504	104.661	308.9071
651.958				
		Total Assets	Total Long-term Debt	EBIT Gross
Profit \	company_name	year		
C_1	1999	740.998	180.447	70.658
191.226	2000	701.854	179.987	45.790
160.444	2001	710.199	217.699	4.711
112.244	2002	686.621	164.658	3.573
109.590	2003	709.292	248.666	20.811
128.656				
		Total Current Liabilities	Retained Earnings \	
company_name	year			
C_1	1999	163.816	201.026	
	2000	125.392	204.065	

2001	150.464	139.603
2002	203.575	124.106
2003	131.261	131.884

		Total Revenue	Total Liabilities	Total Operating Expenses
company_name	year			
C_1	1999	1024.333	401.483	935.302
	2000	874.255	361.642	809.888
	2001	638.721	399.964	611.514
	2002	606.337	391.633	575.592
	2003	651.958	407.608	604.467

info about the dataset:

```
<class 'pandas.core.frame.DataFrame'>
```

MultiIndex: 78682 entries, ('C_1', 1999) to ('C_8971', 2018)

Data columns (total 19 columns):

#	Column	Non-Null Count	Dtype
0	is_bankruptcy	78682 non-null	int64
1	Current assets	78682 non-null	float64
2	Cost of goods sold	78682 non-null	float64
3	Depreciation and amortization	78682 non-null	float64
4	EBITDA	78682 non-null	float64
5	Inventory	78682 non-null	float64
6	Net Income	78682 non-null	float64
7	Total Receivables	78682 non-null	float64
8	Market Value	78682 non-null	float64
9	Net Sales	78682 non-null	float64
10	Total Assets	78682 non-null	float64
11	Total Long-term Debt	78682 non-null	float64
12	EBIT	78682 non-null	float64
13	Gross Profit	78682 non-null	float64
14	Total Current Liabilities	78682 non-null	float64
15	Retained Earnings	78682 non-null	float64
16	Total Revenue	78682 non-null	float64
17	Total Liabilities	78682 non-null	float64
18	Total Operating Expenses	78682 non-null	float64

dtypes: float64(18), int64(1)

memory usage: 12.0+ MB

None

check if there are any NULL value in the dataset:

is_bankruptcy	0
Current assets	0
Cost of goods sold	0
Depreciation and amortization	0
EBITDA	0
Inventory	0
Net Income	0
Total Receivables	0
Market Value	0
Net Sales	0
Total Assets	0
Total Long-term Debt	0
EBIT	0
Gross Profit	0
Total Current Liabilities	0
Retained Earnings	0
Total Revenue	0
Total Liabilities	0
Total Operating Expenses	0

dtype: int64

basic statistics about the dataset (include = "number"):

	is_bankruptcy	Current assets	Cost of goods sold \
count	78682.000000	78682.000000	78682.000000
mean	0.066343	880.362485	1594.529029
std	0.248882	3928.564794	8930.484664
min	0.000000	-7.760000	-366.645000
25%	0.000000	18.924000	17.038250
50%	0.000000	100.449500	103.661000
75%	0.000000	431.526750	634.548000
max	1.000000	169662.000000	374623.000000

	Depreciation and amortization	EBITDA	Inventory \
count	78682.000000	78682.000000	78682.000000
mean	121.234256	376.759424	201.605717
std	652.376804	2012.023142	1060.766096
min	0.000000	-21913.000000	0.000000
25%	1.192000	-0.811000	0.000000
50%	7.929500	15.034500	7.023000
75%	47.971750	139.655250	74.747250
max	28430.000000	81730.000000	62567.000000

	Net Income	Total Receivables	Market Value	Net
Sales \				
count	78682.000000	78682.000000	7.868200e+04	78682.000000

mean	129.382453	286.832743	3.414355e+03	2364.019706
std	1265.532022	1335.978571	1.841410e+04	11950.068842
min	-98696.000000	-0.006000	1.000000e-04	-1964.999000
25%	-7.415750	3.281250	3.498000e+01	27.548500
50%	1.616000	22.820000	2.275118e+02	186.598500
75%	40.144250	131.580500	1.244890e+03	1046.402500
max	104821.000000	65812.000000	1.073391e+06	511729.000000

	Total Assets	Total Long-term Debt	EBIT	Gross Profit \
count	78682.000000	78682.000000	78682.000000	78682.000000
mean	2867.110620	722.483710	255.525035	769.490783
std	12917.944421	3242.170946	1494.643534	3774.703114
min	0.001000	-0.023000	-25913.000000	-
25%	37.363500	0.000000	-2.787000	8.521250
50%	213.203500	7.593500	6.518000	63.581500
75%	1171.364750	248.760750	87.599000	344.074250
max	531864.000000	166250.000000	71230.000000	137106.000000

	Total Current Liabilities	Retained Earnings	Total Revenue \
count	78682.000000	78682.000000	78682.000000
mean	610.072255	532.467069	2364.019706
std	2938.387443	6369.159440	11950.068842
min	0.001000	-102362.000000	-1964.999000
25%	8.889250	-68.282750	27.548500
50%	43.333000	-1.131000	186.598500
75%	222.817000	146.070000	1046.402500
max	116866.000000	402089.000000	511729.000000

	Total Liabilities	Total Operating Expenses
count	78682.000000	78682.000000
mean	1773.563963	1987.260307
std	8053.684902	10419.629038
min	0.001000	-317.197000


25%	13.486000	32.872500
50%	81.988000	168.912000
75%	629.975000	875.522250
max	337980.000000	481580.000000

find the correlation between each numerical variables: shows in heatmap

```
df_correlation = df.drop(['is_bankruptcy'],axis=1)
```

```
plt.figure(figsize=(16,6))
mask = np.triu(np.ones_like(df_correlation.corr(), dtype=bool))
heatmap =
sns.heatmap(df_correlation.corr(),mask=mask,annot=True,cbar=True)
heatmap.set_title("Correlation between each variable",
fontdict={'fontsize': 12})
```

```
Text(0.5, 1.0, 'Correlation between each variable')
```




```
alivedf = df[df['is_bankruptcy']==0]
faileddf = df[df['is_bankruptcy']==1]

year_counts = faileddf.groupby(level='year').size()
print(year_counts)
```

```
plt.figure(figsize=(10, 6))
year_counts.plot(kind='line', xlabel='Year', ylabel='Count',
title='count of failed company each Year')
```

```
plt.xticks(year_counts.index, rotation=90)
plt.show()
```

```
year
1999    380
2000    404
2001    414
2002    414
2003    415
2004    404
2005    379
2006    366
2007    336
2008    284
2009    234
2010    218
2011    194
2012    182
2013    167
2014    142
2015    111
2016     83
2017     57
2018     36
dtype: int64
```



```
print('basic statistics about the alive company dataset (include =
"number"):')
display(alivedf.describe(include = 'number'))
print('basic statistics about the failed company dataset (include =
"number"):')
display(faileddf.describe(include = 'number'))
```

basic statistics about the alive company dataset (include = "number"):

	is_bankruptcy	Current assets	Cost of goods sold \
count	73462.0	73462.000000	73462.000000
mean	0.0	914.542615	1646.982563
std	0.0	4052.047889	9210.587641
min	0.0	-7.760000	-366.645000
25%	0.0	19.235250	17.013500
50%	0.0	102.917500	103.534000
75%	0.0	450.041750	652.468750
max	0.0	169662.000000	374623.000000

	Depreciation and amortization	EBITDA	Inventory \
count	73462.000000	73462.000000	73462.000000
mean	123.746364	393.684586	208.560102
std	671.282060	2076.601504	1089.918086
min	0.000000	-21913.000000	0.000000
25%	1.185250	-0.658000	0.000000

50%	7.935000	15.618000	7.117000
75%	48.001750	144.679000	76.787750
max	28430.000000	81730.000000	62567.000000

	Net Income	Total Receivables	Market Value	Net
Sales \	73462.000000	73462.000000	7.346200e+04	73462.000000
count				
mean	141.994726	297.794547	3.596015e+03	2447.871581
std	1299.259556	1377.396592	1.902253e+04	12326.442146
min	-98696.000000	-0.006000	1.000000e-04	-1964.999000
25%	-6.484750	3.373000	3.590370e+01	27.748000
50%	2.071500	23.454000	2.408953e+02	187.070000
75%	43.444000	136.101500	1.324400e+03	1083.655000
max	104821.000000	65812.000000	1.073391e+06	511729.000000

	Total Assets	Total Long-term Debt	EBIT	Gross
Profit \	73462.000000	73462.000000	73462.000000	
count	73462.000000			
mean	2952.555391	730.516984	269.938080	
800.889131				
std	13290.496814	3309.223955	1542.468732	
3896.084871				
min	0.001000	-0.023000	-25913.000000	-
21536.000000				
25%	37.736750	0.000000	-2.421750	
8.807250				
50%	215.006000	7.092500	7.181000	
64.823000				
75%	1198.552000	245.000000	92.542000	
356.430750				
max	531864.000000	166250.000000	71230.000000	
137106.000000				

	Total Current Liabilities	Retained Earnings	Total Revenue \
count	73462.000000	73462.000000	73462.000000
mean	628.748794	582.034710	2447.871581
std	3021.687870	6570.280016	12326.442146
min	0.001000	-102362.000000	-1964.999000
25%	8.795250	-63.599750	27.748000
50%	43.359500	-0.172500	187.070000
75%	228.800000	158.556000	1083.655000

max	116866.000000	402089.000000	511729.000000
-----	---------------	---------------	---------------

	Total Liabilities	Total Operating Expenses
count	73462.000000	73462.000000
mean	1809.571974	2054.187022
std	8257.442726	10744.450269
min	0.001000	-317.197000
25%	13.318250	32.921250
50%	80.740500	170.290500
75%	634.293250	903.576250
max	337980.000000	481580.000000

basic statistics about the failed company dataset (include = "number"):

	is_bankruptcy	Current assets	Cost of goods sold \
count	5220.0	5220.000000	5220.000000
mean	1.0	399.339353	856.340997
std	0.0	1147.837282	2767.913221
min	1.0	0.001000	-0.666000
25%	1.0	15.315750	17.251500
50%	1.0	75.871500	106.638500
75%	1.0	269.113500	468.845500
max	1.0	16548.000000	40683.000000

	Depreciation and amortization	EBITDA	Inventory	Net
Income \				
count	5220.000000	5220.000000	5220.000000	
mean	85.880919	138.568576	103.735398	-
std	268.526753	521.284649	482.368932	
min	0.000000	-5062.000000	0.000000	-
25%	1.314750	-4.241500	0.076750	-
50%	7.846500	7.790000	6.171000	-
75%	47.442000	85.700000	52.033500	
max	5475.000000	6136.000000	9963.000000	

	Total Receivables	Market Value	Net Sales	Total Assets \
count	5220.000000	5220.000000	5220.000000	5220.000000
mean	132.565307	857.813011	1183.957172	1664.630981
std	421.691639	3396.453790	3567.878694	5284.035224
min	0.000000	0.002000	0.001000	0.015000
25%	2.241750	26.351225	23.961750	33.025000

50%	15.218000	117.799400	180.834500	195.137000
75%	79.708250	495.970825	719.763500	847.652750
max	8207.000000	139092.655000	53012.000000	76995.000000

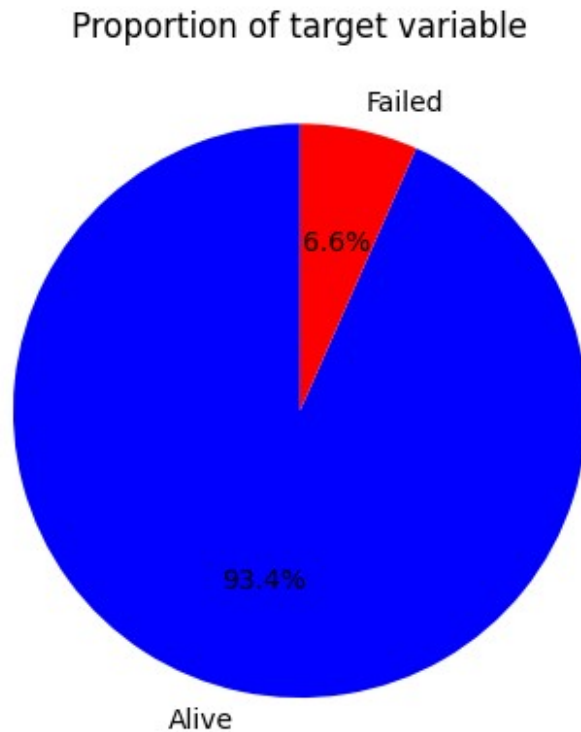
	Total Long-term Debt	EBIT	Gross Profit \
count	5220.000000	5220.000000	5220.000000
mean	609.430007	52.687657	327.616175
std	2077.661309	381.853579	967.555853
min	0.000000	-10537.000000	-4141.334000
25%	0.184000	-9.470250	5.128750
50%	18.438500	0.487000	47.500000
75%	287.792250	37.993500	228.636750
max	21586.000000	4822.000000	15192.000000

	Total Current Liabilities	Retained Earnings	Total Revenue \
count	5220.000000	5220.000000	5220.000000
mean	347.233954	-165.107267	1183.957172
std	1254.398961	1849.943497	3567.878694
min	0.005000	-43091.000000	0.001000
25%	10.015250	-140.766250	23.961750
50%	42.911000	-25.506000	180.834500
75%	160.958750	37.572750	719.763500
max	41695.000000	7832.000000	53012.000000

	Total Liabilities	Total Operating Expenses
count	5220.000000	5220.000000
mean	1266.816748	1045.388596
std	4221.179588	3297.762745
min	0.005000	-0.016000
25%	16.381750	32.143500
50%	97.035000	155.727000
75%	574.765500	617.029000
max	64092.000000	49363.000000

```
import matplotlib.pyplot as plt
import seaborn as sns
bankruptcy_data = {'Alive': 73462, 'Failed': 5220}

plt.pie(bankruptcy_data.values(), labels=bankruptcy_data.keys(),
autopct='%1.1f%%', startangle=90, colors=['blue', 'red'])
plt.title('Proportion of target variable');
```



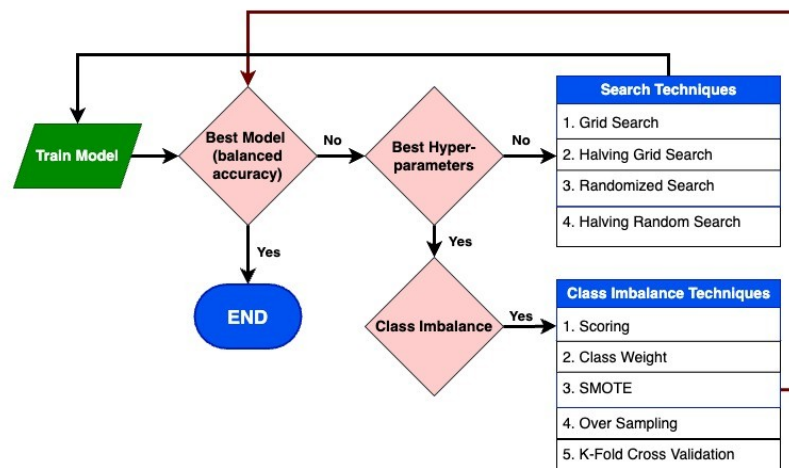
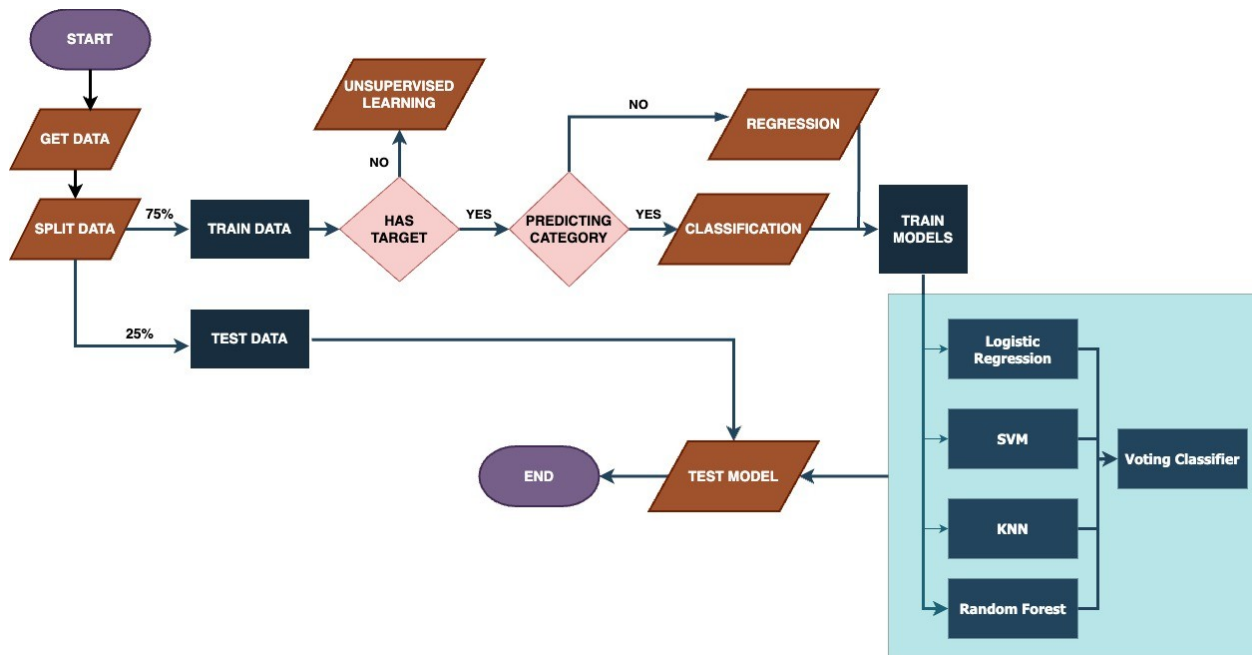
The pie chart presented above unmistakably illustrates a substantial degree of class imbalance within the dataset concerning the target variable. Faced with this challenge, our approach involved the implementation of various techniques to address dataset imbalance comprehensively. Strategies employed included SMOTE (Synthetic Minority Over-sampling Technique), K-Fold Cross-Validation, Bootstrap, and the integration of Balanced Weights.

In our pursuit of mitigating class imbalances, these diverse techniques were systematically applied, each bringing its unique strengths to the endeavor. SMOTE facilitated the generation of synthetic samples for the minority class, K-Fold Cross-Validation enabled robust model evaluation, Bootstrap resampling provided enhanced training data diversity, and Balanced Weights contributed to equitable consideration of minority and majority classes during model training.

##7. Model Selection

Premise: We opt to employ KNN (k Nearest Neighbours), Logistic Regression, Random Forest Classifier, SVM (Support Vector Machine) and Voting Classifier for making the prediction.

We have defined the Model Workflow about our notebook.



Train and Test data split

```

X = df.drop('is_bankruptcy',axis=1)
y = df['is_bankruptcy']

from sklearn.model_selection import train_test_split

```

```
X_train,X_test,y_train,y_test =  
train_test_split(X,y,test_size=0.25,random_state=42)
```

Pipeline

```

from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer, make_column_selector

from sklearn import set_config
set_config(display='diagram') # shows the pipeline graphically when
printed

num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

cat_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('cat_encoder', OneHotEncoder())
])

pipeline = ColumnTransformer([
    ('num', num_pipeline,
     make_column_selector(dtype_include=np.number)),
    ('cat', cat_pipeline,
     make_column_selector(dtype_include='category'))
])

pipeline
ColumnTransformer(transformers=[('num',
                                Pipeline(steps=[('imputer',
                                                    SimpleImputer(strategy='median')),
                                                    ('scaler',
                                                    StandardScaler())])),
                                ('cat',
                                Pipeline(steps=[('imputer',
                                                    SimpleImputer(strategy='most_frequent')),
                                                    ('cat_encoder',
                                                    OneHotEncoder())])),
                                <sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ada0>)],
ColumnTransformer(transformers=[('num',
                                Pipeline(steps=[('imputer',
                                                    SimpleImputer(strategy='median')),
                                                    ('scaler',
                                                    StandardScaler())])),
                                ('cat',
                                Pipeline(steps=[('imputer',
                                                    SimpleImputer(strategy='most_frequent')),
                                                    ('cat_encoder',
                                                    OneHotEncoder())])),
                                <sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ace0>)]])

```

To make it easier for implementation we tried to define functions for the evaluation of the model.

Accuracy and Balanced Accuracy

```
from sklearn.metrics import balanced_accuracy_score, accuracy_score

def calculate_accuracy(y_test, y_pred):
    accuracy = accuracy_score(y_test, y_pred)
    balanced_accuracy = balanced_accuracy_score(y_test, y_pred)
    print(f'Accuracy={accuracy:.4f}, Balanced
Accuracy={balanced_accuracy:.4f}')
```

Confusion Matrix

```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay as CMD

def c_matrix(y_test, y_pred, pipe_line):
    cm = confusion_matrix(y_test, y_pred)
    print("confusion matrix=\n")
    display(cm)
    print()
    # display confusion matrix in chart
    image = CMD(cm, display_labels = pipe_line.classes_)
    image.plot(cmap='coolwarm')
```

Classification Report

```
from sklearn.metrics import classification_report

def class_report(y_test, y_pred):
    class_report = classification_report(y_test, y_pred)
    print('classification report=\n', class_report)
```

ROC Curve

```
from sklearn.metrics import RocCurveDisplay

def roc_curve(pipe_line, X_test, y_test):
    RocCurveDisplay.from_estimator(pipe_line, X_test, y_test)
```

Cost Matrix

```
from sklearn.metrics import make_scorer
from sklearn.model_selection import cross_val_score

def self_define_cost(y_test, y_prediction_binary):
    cm = confusion_matrix(y_test, y_prediction_binary)
```

```

    return cm[0,1]*1 + cm[1,0]*10

def cost_matrix(pipe_line, X_train, y_train):
    custom_cost = make_scorer(self_define_cost)
    cost_simple_pipeline =
cross_val_score(pipe_line,X_train,y_train,scoring=custom_cost,cv=5)
    print('average cost for simple
pipeline=',cost_simple_pipeline.mean())

```

###7.1 Logistic Regression

####Model Training & Evaluation

```

from sklearn.linear_model import LogisticRegression
simple_pipeline = Pipeline([('preprocessor',pipeline),
('classifier',LogisticRegression(random_state=42))])
simple_pipeline

Pipeline(steps=[('preprocessor',
                  ColumnTransformer(transformers=[('num',

Pipeline(steps=[('imputer',
                  SimpleImputer(strategy='median')),
                  ('scaler',
                  StandardScaler())])),
<sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ada0>),
                  ('cat',

Pipeline(steps=[('imputer',
                  SimpleImputer(strategy='most_frequent')),
                  ('cat_encoder',
                  OneHotEncoder())])),
<sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ace0>))]),
                  ('classifier', LogisticRegression(random_state=42))])

simple_pipeline.fit(X_train,y_train) # make a training data

# probaability outcomes
y_prediction_prob = simple_pipeline.predict_proba(X_test)[: , 1] # make
the prediction, the result is shown in probability(y=1)

```

```

print('prediction_probability=',y_prediction_prob)

# change it into binary outcomes
threshold = 0.5 # default
y_prediction_binary = (y_prediction_prob >= threshold).astype(int)
print('prediction_binary=',y_prediction_binary)

# Create a DataFrame with the predicted outcomes and set the index
result = pd.DataFrame({
    'probability': y_prediction_prob,
    'is_bankrupt': y_prediction_binary
}, index=X_test.index)
result

prediction_probability= [0.07176234 0.06008078 0.07089656 ...
0.07501865 0.06558447 0.00627952]
prediction_binary= [0 0 0 ... 0 0 0]

```

company_name	year	probability	is_bankrupt
C_6246	2000	0.071762	0
C_7120	2016	0.060081	0
C_8737	2016	0.070897	0
C_6107	1999	0.072156	0
C_761	2002	0.073810	0
...	
C_4035	2004	0.070763	0
C_1651	2009	0.073674	0
C_8176	2008	0.075019	0
C_2263	2010	0.065584	0
C_3596	2010	0.006280	0

[19671 rows x 2 columns]

#####Balanced Accuracy Score

```

calculate_accuracy(y_test, y_prediction_binary)

Accuracy=0.9320, Balanced Accuracy=0.5031

```

#####Confusion Matrix


```

c_matrix(y_test, y_prediction_binary, simple_pipeline)

confusion matrix=

array([[18324,    8],
       [ 1330,    9]])

```



#####Classification Report


```
class_report(y_test, y_prediction_binary)
```

```
classification report=
```

	precision	recall	f1-score	support
0	0.93	1.00	0.96	18332
1	0.53	0.01	0.01	1339
accuracy			0.93	19671
macro avg	0.73	0.50	0.49	19671
weighted avg	0.90	0.93	0.90	19671

#####ROC Curve

```
roc_curve(simple_pipeline, X_test, y_test)
```

#####Cost Matrix

```
cost_matrix(simple_pipeline,X_train,y_train)
average cost for simple pipeline= 7709.2
```

####Parameter Tuning

#####GridSearchCV

```
from sklearn.model_selection import GridSearchCV

param_grid= {'LogisticRegression__penalty':['l1',
'l2'],'LogisticRegression__C':[0.1, 1.0,
10.0],'LogisticRegression__class_weight': [None, 'balanced']}

grid_search = GridSearchCV(estimator=simple_pipeline,
param_grid=param_grid,cv=5, scoring='balanced_accuracy')
grid_search.fit(X_train,y_train)

grid_search_all_result = pd.DataFrame(grid_search.cv_results_)
#display(grid_search_all_result)
result =
grid_search_all_result.sort_values(by='mean_test_score',ascending=False)
e)
```

```
display(result.filter(regex= '^(^param_|mean_test_score)',
axis=1).head())
```

```
    # regular expression
    # ^ : start from
    # | : or
```

```
print()
print('The best parameters are=', grid_search.best_params_)
```

	param_LogisticRegression__C	param_LogisticRegression__class_weight
3	0.1	balanced
7	1.0	balanced
11	10.0	balanced
9	10.0	None
5	1.0	None

	param_LogisticRegression__penalty	mean_test_score
3	l2	0.594286
7	l2	0.594026
11	l2	0.593970
9	l2	0.503778
5	l2	0.503538

```
The best parameters are= {'LogisticRegression__C': 0.1,
'LogisticRegression__class_weight': 'balanced',
'LogisticRegression__penalty': 'l2'}
```

#####HalvingGridSearchCV

```
from sklearn.experimental import enable_halving_search_cv
from sklearn.model_selection import HalvingGridSearchCV

param_grid= {'LogisticRegression__penalty':['l1',
'l2'],'LogisticRegression__C':[0.1, 1.0,
10.0],'LogisticRegression__class_weight': [None, 'balanced']}

halving_grid_search =
HalvingGridSearchCV(estimator=simple_pipeline,param_grid=param_grid,cv
=5, scoring='balanced_accuracy',min_resources='exhaust')
halving_grid_search.fit(X_train,y_train)

halving_grid_search_all_result =
pd.DataFrame(halving_grid_search.cv_results_)
#display(halving_grid_search_all_result)
```

```

result =
halving_grid_search_all_result.sort_values(by=['iter','mean_test_score'],ascending=False)
display(result.filter(regex='^param_|mean_test_score',axis=1).head())
# regular expression
# ^ : start from
# | : or

print()
print('The best parameters are=', halving_grid_search.best_params_)

```

	param_LogisticRegression__C	param_LogisticRegression__class_weight
16	10.0	balanced
17	1.0	balanced
13	1.0	balanced
15	10.0	balanced
14	0.1	balanced

	param_LogisticRegression__penalty	mean_test_score
16	l2	0.594007
17	l2	0.594000
13	l2	0.595186
15	l2	0.594675
14	l2	0.592788

```

The best parameters are= {'LogisticRegression__C': 10.0,
'LogisticRegression__class_weight': 'balanced',
'LogisticRegression__penalty': 'l2'}

```

#####RandomSearchCV

```

from sklearn.model_selection import RandomizedSearchCV

param_distributions= {'LogisticRegression__penalty':['l1',
'l2'],'LogisticRegression__C':[0.1, 1.0,
10.0],'LogisticRegression__class_weight': [None, 'balanced']}

random_search =
RandomizedSearchCV(estimator=simple_pipeline,param_distributions=param
_distributions,n_iter=20,cv=5,scoring='balanced_accuracy',random_state=42)
random_search.fit(X_train,y_train)

random_search_all_result = pd.DataFrame(random_search.cv_results_)
#display(random_search_all_result)

```

```

result =
random_search_all_result.sort_values(by='mean_test_score',ascending=False)
display(result.filter(regex = '^param_|
mean_test_score',axis=1).head())
print()
print('The best parameters are=', random_search.best_params_)

```

```

    param_LogisticRegression__penalty
param_LogisticRegression__class_weight \
3                                l2
balanced
7                                l2
balanced
11                               l2
balanced
9                                l2
None
5                                l2
None

```

	param_LogisticRegression__C	mean_test_score
3	0.1	0.594286
7	1.0	0.594026
11	10.0	0.593970
9	10.0	0.503778
5	1.0	0.503538

```

The best parameters are= {'LogisticRegression__penalty': 'l2',
'LogisticRegression__class_weight': 'balanced',
'LogisticRegression__C': 0.1}

```

#####HalvingRandomSearchCV

```

# halving random search
from sklearn.experimental import enable_halving_search_cv
from sklearn.model_selection import HalvingRandomSearchCV

param_distribs= {'LogisticRegression__penalty':['l1',
'l2'],'LogisticRegression__C':[0.1, 1.0,
10.0],'LogisticRegression__class_weight': [None, 'balanced']}

halving_random_search =
HalvingRandomSearchCV(estimator=simple_pipeline,param_distributions=pa
ram_distribs,n_candidates=20,cv=5,scoring='balanced_accuracy',random_s
tate=42,min_resources='exhaust')
halving_random_search.fit(X_train,y_train)

halving_random_search_all_result =

```

```
pd.DataFrame(halving_random_search.cv_results_)
#display(halving_random_search_all_result)
result =
halving_random_search_all_result.sort_values(by=['iter','mean_test_score'],ascending=False)
display(result.filter(regex='^param|mean_test_score',axis=1).head())
print()
print('The best parameters are=', halving_random_search.best_params_)
```

```
    param_LogisticRegression__penalty
param_LogisticRegression__class_weight \
16                                     l2
balanced
17                                     l2
balanced
14                                     l2
balanced
13                                     l2
balanced
15                                     l2
balanced
```

```
    param_LogisticRegression__C \
16                               0.1
17                              10.0
14                              10.0
13                               0.1
15                               1.0
```

	params	mean_test_score
16	<code>{'LogisticRegression__penalty': 'l2', 'Logisti...</code>	0.594235
17	<code>{'LogisticRegression__penalty': 'l2', 'Logisti...</code>	0.593991
14	<code>{'LogisticRegression__penalty': 'l2', 'Logisti...</code>	0.589938
13	<code>{'LogisticRegression__penalty': 'l2', 'Logisti...</code>	0.589215
15	<code>{'LogisticRegression__penalty': 'l2', 'Logisti...</code>	0.588895

```
The best parameters are= {'LogisticRegression__penalty': 'l2',
'LogisticRegression__class_weight': 'balanced',
'LogisticRegression__C': 0.1}
```

####Model Retraining

```
# after find the best parameters, re-predict the X_test
advanced_pipeline = Pipeline([
```

```

    ('preprocessor', pipeline),
    ('classifier', LogisticRegression(random_state=42, penalty='l2',
class_weight='balanced', C=0.1))])
advanced_pipeline.fit(X_train,y_train)

# probaability outcomes
y_prediction_prob = advanced_pipeline.predict_proba(X_test)[:, 1]
print('prediction_probability=',y_prediction_prob)

# change it into binary outcomes
threshold = 0.5 # default
y_prediction_binary = (y_prediction_prob >= threshold).astype(int)
print('prediction_binary=',y_prediction_binary)

# Create a DataFrame with the predicted outcomes and set the index
result = pd.DataFrame({
    'probability': y_prediction_prob,
    'is_bankrupt': y_prediction_binary
}, index=X_test.index)
result

prediction_probability= [0.56034198 0.47455307 0.51861262 ...
0.52252278 0.48669172 0.06608786]
prediction_binary= [1 0 1 ... 1 0 0]

```

company_name	year	probability	is_bankrupt
C_6246	2000	0.560342	1
C_7120	2016	0.474553	0
C_8737	2016	0.518613	1
C_6107	1999	0.522009	1
C_761	2002	0.518097	1
...
C_4035	2004	0.532785	1
C_1651	2009	0.518733	1
C_8176	2008	0.522523	1
C_2263	2010	0.486692	0
C_3596	2010	0.066088	0

[19671 rows x 2 columns]

####Model Revaluation

#####Balanced Accuracy Score

```


calculate_accuracy(y_test,y_prediction_binary)

Accuracy=0.3629, Balanced Accuracy=0.5945

```

#####Confusion Matrix

```
c_matrix(y_test, y_prediction_binary, advanced_pipeline)
confusion matrix=
array([[ 5984, 12348],
       [  184, 1155]])
```




#####Classification Report

```
class_report(y_test, y_prediction_binary)
classification report=
```

	precision	recall	f1-score	support
0	0.97	0.33	0.49	18332
1	0.09	0.86	0.16	1339
accuracy			0.36	19671
macro avg	0.53	0.59	0.32	19671
weighted avg	0.91	0.36	0.47	19671

#####ROC Curve

```
roc_curve(advanced_pipeline, X_test, y_test)
```



#####Cost Matrix

```
cost_matrix(advanced_pipeline,X_train,y_train)  
average cost for simple pipeline= 8480.0
```

#####Techniques for Dataset Imbalance

Premise: Enhance the advanced_pipeline to assess accuracy, enabling improved predictions on unseen data using k-fold and bootstrap resampling.

#####K-fold

```
# k-fold  
from sklearn.model_selection import cross_val_score  
cv = cross_val_score(advanced_pipeline,  
X_train,y_train,cv=5,scoring='balanced_accuracy')  
print('accuracy after using cv=',cv.mean())  
accuracy after using cv= 0.5942864150883247
```


#####Bootstrap

```
# bootstrap
from sklearn.utils import resample
from sklearn.metrics import balanced_accuracy_score
import numpy as np

def bootstrap(model, X_train, y_train, X_test, y_test,
              n_iterations=100):
    accuracy_scores = []

    for times in range(n_iterations):
        # Bootstrap resampling on the minority class
        X_resample, y_resample = resample(X_train[y_train == 1],
                                          y_train[y_train == 1], replace=True, n_samples=X_train[y_train ==
0].shape[0])
        X_resample = np.concatenate((X_resample, X_train[y_train ==
0]), axis=0) # Combine minority and majority class along axis 0
        y_resample = np.concatenate((y_resample, y_train[y_train ==
0]), axis=0)
        X_resample = pd.DataFrame(X_resample, columns=['Current
assets',
'Cost of goods sold', 'Depreciation and amortization',
'EBITDA',
'Inventory', 'Net Income', 'Total Receivables', 'Market Value',
'Net Sales', 'Total Assets', 'Total Long-term Debt', 'EBIT',
'Gross Profit', 'Total Current Liabilities', 'Retained
Earnings',
'Total Revenue', 'Total Liabilities', 'Total Operating
Expenses'])
        y_resample = pd.DataFrame(y_resample,
columns=['is_bankruptcy'])

        model.fit(X_resample, y_resample)
        y_pred = model.predict(X_test)
        accuracy = balanced_accuracy_score(y_test, y_pred)
        accuracy_scores.append(accuracy)

    return np.array(accuracy_scores)

bootstrap_scores = bootstrap(advanced_pipeline, X_train, y_train,
X_test, y_test)
print("Balanced accuracy after using bootstrap:",
bootstrap_scores.mean())
```

Balanced accuracy after using bootstrap: 0.5953791681828338

#####SMOTE Technique

```

#smote
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline, make_pipeline
from sklearn.metrics import accuracy_score, balanced_accuracy_score

smote = SMOTE(random_state=42)

# Create the pipeline
logis_smote_pipe = Pipeline([
    ('transformer', pipeline),
    ('smote', smote),
    ('LogisticRegression', LogisticRegression(random_state=42,
penalty='l2', class_weight='balanced', C=0.1))
])

logis_smote_pipe.fit(X_train, y_train)
y_pred2_smote = logis_smote_pipe.predict(X_test)
accuracy = accuracy_score(y_test, y_pred2_smote)
balanced_accuracy = balanced_accuracy_score(y_test, y_pred2_smote)
print(f'Accuracy={accuracy:.4f}, Balanced
Accuracy={balanced_accuracy:.4f}')

Accuracy=0.3701, Balanced Accuracy=0.5946

```

Following the training and prediction phase of our Logistic Regression model, we identified the optimal hyperparameters using the Halving Random Search method. The hyperparameters that yielded the best performance were found to be `penalty='l2'`, `class_weight='balanced'`, and `C=0.1`.

Incorporating bootstrapping techniques, our model achieved a balanced accuracy of 0.5953. This signifies a balanced and robust predictive capability, especially crucial in scenarios where the dataset is imbalanced.

The adoption of Halving Random Search not only facilitated the discovery of hyperparameters that enhance model performance but also underscored the effectiveness of an intelligent search strategy in navigating the hyperparameter space. This comprehensive approach ensures that our Logistic Regression model is fine-tuned for optimal performance, laying the foundation for reliable predictions in the face of imbalanced data.

###7.2 Random Forest

####Model Training

```

#create pipeline
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(random_state=42)

rf_pipeline = Pipeline([
    ('preprocessor', pipeline),

```

```

    ('model', rf)
])

#fit a RandomForestClassifier classifier
rf_pipeline.fit(X_train, y_train)

Pipeline(steps=[('preprocessor',
                  ColumnTransformer(transformers=[('num',

Pipeline(steps=[('imputer',
                  SimpleImputer(strategy='median')),

                  ('scaler',
                  StandardScaler())])),

<sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ada0>),

                  ('cat',

Pipeline(steps=[('imputer',
                  SimpleImputer(strategy='most_frequent')),

                  ('cat_encoder',
                  OneHotEncoder())])),

<sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ace0>)])),

    ('model', RandomForestClassifier(random_state=42))])

transformed_data = pipeline.fit_transform(df[['Current assets', 'Cost
of goods sold', 'Depreciation and
amortization', 'EBITDA', 'Inventory', 'Net Income', 'Total Receivables',
        'Market Value', 'Net Sales', 'Total Assets', 'Total Long-term
Debt', 'EBIT', 'Gross Profit', 'Total Current Liabilities', 'Retained
Earnings',
        'Total Revenue', 'Total Liabilities', 'Total Operating
Expenses']])
transformed_df = pd.DataFrame(transformed_data,
columns=pipeline.get_feature_names_out())
transformed_df

```

	num__Current assets	num__Cost of goods sold \
0	-0.093952	-0.085262
1	-0.100421	-0.098620
2	-0.112944	-0.119597
3	-0.123188	-0.122926
4	-0.114078	-0.119952

...
78677	-0.164731	-0.173697
78678	-0.197224	-0.171923
78679	-0.207810	-0.170815
78680	-0.189677	-0.171101
78681	-0.203071	-0.170844

	num__Depreciation and amortization	num__EBITDA	num__Inventory
\			
0	-0.157673	-0.143005	0.126713
1	-0.157360	-0.155264	0.112169
2	-0.151353	-0.173733	0.080115
3	-0.144185	-0.171974	0.055006
4	-0.144939	-0.163651	0.043025
...
78677	-0.164232	-0.164584	-0.186875
78678	-0.120548	-0.086792	-0.187901
78679	-0.086112	-0.147966	-0.187625
78680	-0.085694	-0.152876	-0.188160
78681	-0.085891	-0.153825	-0.188067

	num__Net Income	num__Total Receivables	num__Market Value	\
0	-0.074451	-0.118629	-0.165179	
1	-0.087593	-0.128480	-0.164942	
2	-0.148809	-0.156669	-0.165622	
3	-0.112042	-0.165057	-0.177638	
4	-0.099467	-0.136359	-0.168646	
...	
78677	-0.082275	-0.197599	-0.144340	
78678	0.000241	-0.173823	-0.156771	
78679	-0.103376	-0.182913	-0.153984	
78680	-0.118357	-0.194328	-0.163014	
78681	-0.142493	-0.180389	-0.166188	

	num__Net Sales	num__Total Assets	num__Total Long-term Debt	\
0	-0.112108	-0.164587	-0.167184	
1	-0.124667	-0.167617	-0.167326	
2	-0.144377	-0.166971	-0.155694	
3	-0.147087	-0.168796	-0.172054	
4	-0.143269	-0.167041	-0.146143	

78677	-0.189104	-0.136865	-0.165883
78678	-0.173462	-0.077504	0.014688
78679	-0.183612	-0.086769	-0.011874
78680	-0.184279	-0.087553	-0.008775
78681	-0.184394	-0.096126	-0.027871

	num__EBIT	num__Gross Profit	num__Total Current Liabilities \
0	-0.123687	-0.153196	-0.151872
1	-0.140325	-0.161351	-0.164949
2	-0.167810	-0.174120	-0.156416
3	-0.168571	-0.174823	-0.138341
4	-0.157038	-0.169772	-0.162951
78677	-0.149872	-0.187726	-0.198027
78678	-0.064219	-0.142402	-0.177631
78679	-0.161599	-0.177156	-0.178435
78680	-0.168392	-0.178594	-0.179713
78681	-0.169583	-0.179564	-0.180613

	num__Retained Earnings	num__Total Revenue	num__Total Liabilities \
0	-0.052039	-0.112108	-
0.170368			
1	-0.051562	-0.124667	-
0.175315			
2	-0.061683	-0.144377	-
0.170557			
3	-0.064116	-0.147087	-
0.171591			
4	-0.062895	-0.143269	-
0.169607			
...	
78677	-0.079190	-0.189104	-
0.192171			
78678	-0.058828	-0.173462	-
0.110911			
78679	-0.059055	-0.183612	-
0.124581			
78680	-0.062258	-0.184279	-
0.123779			
78681	-0.070257	-0.184394	-
0.131726			

	num__Total Operating Expenses
0	-0.100960
1	-0.112996
2	-0.132035
3	-0.135482

```

4 -0.132711
...
78677 -0.185099
78678 -0.182180
78679 -0.182009
78680 -0.181826
78681 -0.181774

```

```
[78682 rows x 18 columns]
```

####Model Evaluation

#####Balanced Accuracy Score

```

y_pred_prob = rf_pipeline.predict_proba(X_test)[: , 1]
y_pred = (y_pred_prob > 0.5).astype(int)
calculate_accuracy(y_test, y_pred)

```

```
Accuracy=0.9343, Balanced Accuracy=0.5217
```

#####Confusion Matrix


```
c_matrix(y_test, y_pred, rf_pipeline)
```

```
confusion matrix=
```

```

array([[18320, 12],
       [ 1280, 59]])

```



#####Classification Report


```
class_report(y_test, y_pred)
```

```
classification report=
```

	precision	recall	f1-score	support
0	0.93	1.00	0.97	18332
1	0.83	0.04	0.08	1339
accuracy			0.93	19671
macro avg	0.88	0.52	0.52	19671
weighted avg	0.93	0.93	0.91	19671

#####ROC Curve

```
roc_curve(rf_pipeline, X_test, y_test)
```




#####Cost Matrix

```
cost_matrix(rf_pipeline, X_train, y_train)
average cost for simple pipeline= 7304.0
```

#####Feature importance

```
import pandas as pd
import matplotlib.pyplot as plt
# Create a pd.Series of features importances
importances_rf = pd.Series(rf.feature_importances_, index = X.columns)
# Sort importances_rf
sorted_importances_rf = importances_rf.sort_values()
# Make a horizontal bar plot
sorted_importances_rf.plot(kind='barh', color='darkred'); plt.show()
```

####Parameter Tuning

#####GridSearchCV

```

grid_max_depth = [None,2,5,10]
grid_min_samples_leaf = [1, 5, 10]
n_estimators = [100,200,300]

param_grid_g = [{'model__max_depth': grid_max_depth,
                  'model__min_samples_leaf':
grid_min_samples_leaf , 'model__n_estimators':n_estimators} ]

from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(rf_pipeline, param_grid_g,
scoring='balanced_accuracy')
grid_search.fit(X_train, y_train)
print('The best parameters are ', grid_search.best_params_)

The best parameters are {'model__max_depth': None,
 'model__min_samples_leaf': 1, 'model__n_estimators': 100}

#show top5 parameters combinations
grid_cv_result = pd.DataFrame(grid_search.cv_results_)
grid_cv_result.sort_values(by='mean_test_score', ascending=False,
inplace=True)
grid_cv_result.filter(regex = '^param_|mean_test_score',
axis=1).head()

```

	param_model__max_depth	param_model__min_samples_leaf	\
0	None		1
1	None		1
2	None		1
3	None		5
4	None		5

	param_model__n_estimators	mean_test_score
0	100	0.529578
1	200	0.529264
2	300	0.528988
3	100	0.508430
4	200	0.507933

#####HalvingGridSearchCV

```
hgrid_max_depth = [None,2,5,10,12]
hgrid_min_samples_leaf = [1,5,10,15]
n_estimators = [100,200,300]
```

```
param_grid_hg = [{'model__max_depth': hgrid_max_depth,
                  'model__min_samples_leaf':
hgrid_min_samples_leaf, 'model__n_estimators':n_estimators} ]
```

```
from sklearn.experimental import enable_halving_search_cv
from sklearn.model_selection import HalvingGridSearchCV
```

```
hgrid_search = HalvingGridSearchCV(rf_pipeline, param_grid_hg,
scoring='balanced_accuracy',random_state=42)
hgrid_search.fit(X_train, y_train)
print('The best parameters are ', hgrid_search.best_params_)
```

```
The best parameters are {'model__max_depth': None,
'model__min_samples_leaf': 1, 'model__n_estimators': 100}
```

```
hgrid_cv_results = pd.DataFrame(hgrid_search.cv_results_)
hgrid_cv_results.sort_values(by=['iter','mean_test_score'],
ascending=False, inplace=True)
hgrid_cv_results.filter(regex = '(iter|^param_|mean_test_score)',
axis=1).head()
```

	iter	param_model__max_depth	param_model__min_samples_leaf	\
89	3	None		1
88	3	None		1
87	3	None		1
86	2	None		1
85	2	None		1

	param_model__n_estimators	mean_test_score
89	100	0.529603

88	300	0.529371
87	200	0.529362
86	100	0.513943
85	300	0.513207

#####RandomSearchCV

```

random_max_depth = [None,25,40,55]
random_min_samples_leaf = list(range(1,20,5))

param_grid_r = [{'model__max_depth': random_max_depth,
                  'model__min_samples_leaf':
random_min_samples_leaf, 'model__n_estimators':n_estimators} ]

from sklearn.model_selection import RandomizedSearchCV

random_search = RandomizedSearchCV(rf_pipeline, param_grid_r,
scoring='balanced_accuracy',random_state=42)
random_search.fit(X_train, y_train)
print('The best parameters are ', random_search.best_params_)

The best parameters are {'model__n_estimators': 100,
'model__min_samples_leaf': 1, 'model__max_depth': 25}

#show top5 parameters combinations
random_cv_result = pd.DataFrame(random_search.cv_results_)
random_cv_result.sort_values(by='mean_test_score', ascending=False,
inplace=True)
random_cv_result.filter(regex = '^(^param_|mean_test_score)',
axis=1).head()

```

	param_model__n_estimators	param_model__min_samples_leaf	\
6	100		1
9	200		1
2	300		1
5	200		1
4	100		1

	param_model__max_depth	mean_test_score
6	25	0.529697
9	40	0.529530
2	40	0.529402
5	55	0.529264
4	40	0.529182

#####HalvingRandomSearchCV

```

#HalvingRandomSearchCV
halving_max_depth = [None,30,50,70]

```

```

halving_min_samples_leaf = list(range(1,20,5))

param_grid_h = {'model__max_depth': halving_max_depth,
                'model__min_samples_leaf':
                halving_min_samples_leaf, 'model__n_estimators': n_estimators}

from sklearn.experimental import enable_halving_search_cv
from sklearn.model_selection import HalvingRandomSearchCV

halving_random_search = HalvingRandomSearchCV(rf_pipeline,
param_grid_h, scoring='balanced_accuracy', random_state=42)
halving_random_search.fit(X_train, y_train)
print('The best parameters are ', halving_random_search.best_params_)

```

The best parameters are {'model__n_estimators': 200, 'model__min_samples_leaf': 16, 'model__max_depth': 70}

```

#show top5 parameters combinations
#incorporate 'iter' into sorting values(using most data)
hrandom_cv_results = pd.DataFrame(halving_random_search.cv_results_)
hrandom_cv_results.sort_values(by=['iter', 'mean_test_score'],
ascending=False, inplace=True)
hrandom_cv_results.filter(regex = '(iter|^param_|mean_test_score)',
axis=1).head()

```

	iter	param_model__n_estimators	param_model__min_samples_leaf	\
70	3	200	16	
71	3	300	16	
64	2	100	11	
65	2	200	11	
66	2	300	11	

	param_model__max_depth	mean_test_score
70	70	0.5
71	70	0.5
64	30	0.5
65	30	0.5
66	30	0.5

####Model Retraining

```

rf2 = RandomForestClassifier(random_state=42, min_samples_leaf= 1,
max_depth = 25, n_estimators = 100)

rf_pipeline2 = Pipeline([
    ('preprocessor', pipeline),
    ('model', rf2)
])

```

```

#fit a RandomForestClassifier classifier
rf_pipeline2.fit(X_train, y_train)

Pipeline(steps=[('preprocessor',
                  ColumnTransformer(transformers=[('num',
Pipeline(steps=[('imputer',
                  SimpleImputer(strategy='median')),
                  ('scaler',
                  StandardScaler())])),
                  <sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ada0>),
                  ('cat',
                  Pipeline(steps=[('imputer',
                  SimpleImputer(strategy='most_frequent')),
                  ('cat_encoder',
                  OneHotEncoder())])),
                  <sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ace0>)])),
          ('model',
          RandomForestClassifier(max_depth=25,
random_state=42))])

```

####Model Reevaluation

#####Balanced Accuracy Score

```

y_pred_prob2 = rf_pipeline2.predict_proba(X_test)[: , 1]
y_pred2 = (y_pred_prob2 > 0.5).astype(int)
calculate_accuracy(y_test, y_pred2)

Accuracy=0.9355, Balanced Accuracy=0.5272

```

#####Confusion Matrix


```

c_matrix(y_test, y_pred2, rf_pipeline2)

confusion matrix=

array([[18330,    2],
       [ 1266,   73]])

```



#####Classification Report


```
class_report(y_test, y_pred2)
```

```
classification report=
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	18332
1	0.97	0.05	0.10	1339
accuracy			0.94	19671
macro avg	0.95	0.53	0.53	19671
weighted avg	0.94	0.94	0.91	19671

#####ROC Curve

```
roc_curve(rf_pipeline2, X_test, y_test)
```



#####Cost Matrix

```
cost_matrix(rf_pipeline2, X_train, y_train)
average cost for simple pipeline= 7302.2
```

####Techniques for Dataset Imbalance

#####K-fold

```
#k-fold
from sklearn.model_selection import cross_val_score
cv2 =
cross_val_score(rf_pipeline2,X_train,y_train,cv=5,scoring='balanced_ac
curacy')

print(f"Mean Balanced Accuracy: {np.mean(cv2):.4f}")
Mean Balanced Accuracy: 0.5297
```

#####SMOTE Technique

```
#smote
from imblearn.over_sampling import SMOTE
```

```

from imblearn.pipeline import Pipeline, make_pipeline
smote = SMOTE(random_state=42)

# Create the pipeline
rf_smote_pipe = Pipeline([
    ('preprocessor', pipeline),
    ('smote', smote),
    ('model', rf2)
])

rf_smote_pipe.fit(X_train, y_train)
y_pred2_smote = rf_smote_pipe.predict(X_test)
accuracy = accuracy_score(y_test, y_pred2_smote)
balanced_accuracy = balanced_accuracy_score(y_test, y_pred2_smote)
print(f'Accuracy={accuracy:.4f}, Balanced
Accuracy={balanced_accuracy:.4f}')

Accuracy=0.9123, Balanced Accuracy=0.7055

```

The initial implementation of the Random Forest model revealed a suboptimal balanced accuracy of 0.5217, prompting the adoption of a Random Search approach. This exploration led to the discovery of key hyperparameters—`min_samples_leaf=1`, `max_depth=25`, and `n_estimators=100`—that significantly improved model performance.

Upon retraining the model with these optimized hyperparameters and implementing strategies to tackle class imbalances, notably leveraging the SMOTE technique, we observed a substantial increase in balanced accuracy to a noteworthy level of 0.7055. This enhancement underscores the critical role of hyperparameter tuning and targeted handling of class imbalances, collectively contributing to the model's heightened predictive capabilities.

Notably, this refinement in the Random Forest model coincided with a remarkable increase in the AUC value, soaring from an initial 0.78 to an impressive 0.85. This noteworthy improvement in the AUC metric further accentuates the efficacy of our iterative and systematic approach. The model, now equipped with optimal hyperparameters and adept at addressing imbalanced data, stands as a testament to our commitment to refining its performance for robust and reliable predictions.

###7.3 KNN

####Model Training

```

from sklearn.pipeline import Pipeline
from sklearn.neighbors import KNeighborsClassifier

knn_pipeline = Pipeline([
    ('preprocessing', pipeline),
    ('knn', KNeighborsClassifier())
])

knn_pipeline

```



```

Pipeline(steps=[('preprocessing',
                  ColumnTransformer(transformers=[('num',
Pipeline(steps=[('imputer',
SimpleImputer(strategy='median'))],
('scaler',
StandardScaler())])),
<sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ada0>),
('cat',
Pipeline(steps=[('imputer',
SimpleImputer(strategy='most_frequent'))],
('cat_encoder',
OneHotEncoder()))]),
<sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ace0>)])),
('knn', KNeighborsClassifier())])

```

####Model Evaluation

```

knn_pipeline.fit(X_train, y_train)
y_pred = knn_pipeline.predict(X_test)

```

#####Balanced Accuracy Score

```

calculate_accuracy(y_test, y_pred)
Accuracy=0.9385, Balanced Accuracy=0.6029

```


#####Confusion Matrix

```

c_matrix(y_test, y_pred, knn_pipeline)
confusion matrix=

array([[18174, 158],
       [ 1052, 287]])

```



#####Classification Report

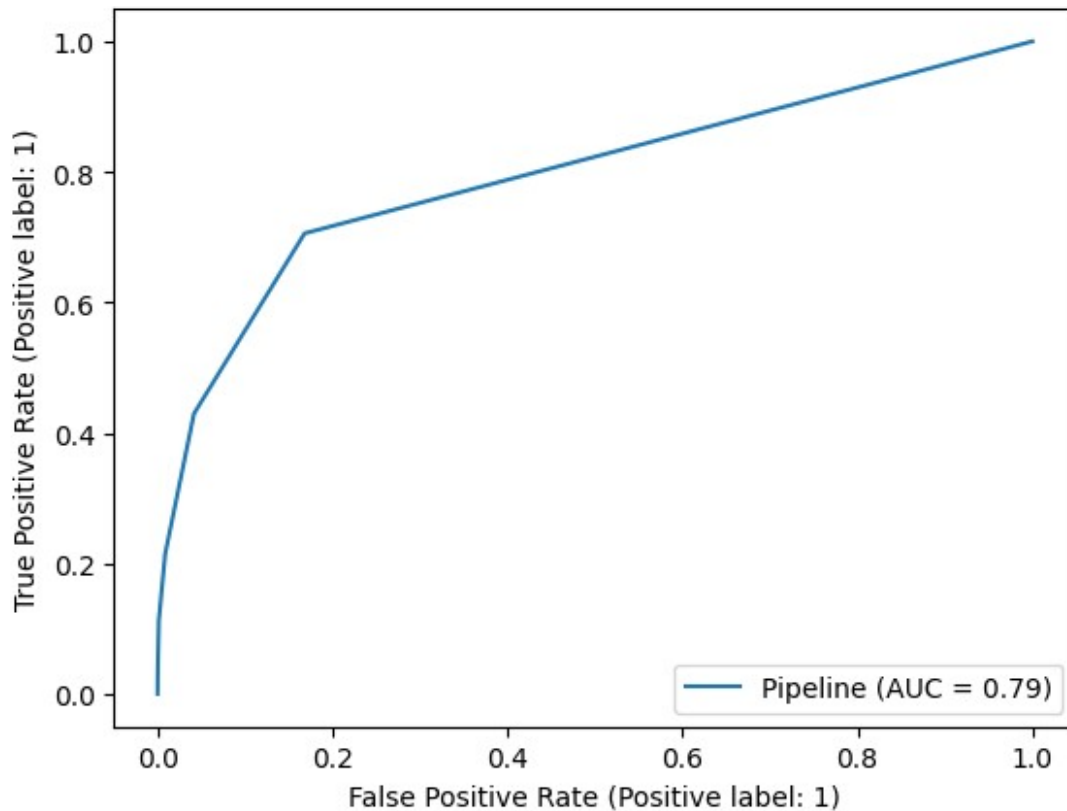
```
class_report(y_test, y_pred)
```

```
classification report=
```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	18332
1	0.64	0.21	0.32	1339
accuracy			0.94	19671
macro avg	0.80	0.60	0.64	19671
weighted avg	0.92	0.94	0.92	19671

#####ROC Curve

```
roc_curve(knn_pipeline, X_test, y_test)
```



#####Cost Matrix

```
cost_matrix(knn_pipeline, X_train, y_train)
average cost for simple pipeline= 6379.4
```

####Parameter Tuning

#####GridSearchCV

```
from sklearn.model_selection import GridSearchCV
param_grid = [
    {'knn__n_neighbors': np.arange(2, 11, 2),
     'knn__p': [1, 2]
    },
]
# Check what's in this parameter grid
print('The parameter grid : ')
print(param_grid)

grid_search = GridSearchCV(knn_pipeline, param_grid, cv=3,
scoring='balanced_accuracy')
grid_search.fit(X_train, y_train)
print('\n\nThe best parameters are ', grid_search.best_params_)
```

```
grid_cv_res = pd.DataFrame(grid_search.cv_results_) # convert to DF for convenience
grid_cv_res.sort_values(by='mean_test_score', ascending=False, inplace=True)
grid_cv_res.filter(regex = '^param_|mean_test_score', axis=1).head()
```

The parameter grid :

```
[{'knn__n_neighbors': array([ 2,  4,  6,  8, 10]), 'knn__p': [1, 2]}]
```

The best parameters are {'knn__n_neighbors': 2, 'knn__p': 1}

	param_knn__n_neighbors	param_knn__p	mean_test_score
0	2	1	0.603385
1	2	2	0.592317
2	4	1	0.573312
3	4	2	0.562930
4	6	1	0.550195

#####HalvingGridSearchCV

```
from sklearn.experimental import enable_halving_search_cv
from sklearn.model_selection import HalvingGridSearchCV

param_grid = [
    {'knn__n_neighbors': np.arange(2, 11, 2),
     'knn__p': [1, 2]
    },
]
# Check what's in this parameter grid
print('The parameter grid : ')
print(param_grid)

# Change to new strategy starting here
halving_grid_search = HalvingGridSearchCV(knn_pipeline, param_grid,
cv=3,
                                         min_resources='exhaust', # use all
data in the last round, start with as needed
                                         scoring='balanced_accuracy')
halving_grid_search.fit(X_train, y_train)
print('The best parameters are ', halving_grid_search.best_params_)

halving_grid_cv_res = pd.DataFrame(halving_grid_search.cv_results_) # convert to DF for convenience
# In the end, we care about performances in the last iteration (using most data)
# So, let's sort by iteration (descending), then by test score (descending)
halving_grid_cv_res.sort_values(by=['iter', 'mean_test_score'], ascending=False, inplace=True)
```

```
# and check the top few rows
halving_grid_cv_res.filter(regex = '(iter|^param_|mean_test_score|
n_resources)', axis=1).head()
```

The parameter grid :

```
[{'knn__n_neighbors': array([ 2,  4,  6,  8, 10]), 'knn__p': [1, 2]}]
```

The best parameters are {'knn__n_neighbors': 2, 'knn__p': 1}

	iter	n_resources	param_knn__n_neighbors	param_knn__p	mean_test_score
15	2	59004	2	1	0.603410
14	2	59004	2	2	0.592340
13	1	19668	2	1	0.542865
12	1	19668	2	2	0.536418
11	1	19668	4	1	0.521296

#####RandomSearchCV

```
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

param_distributions = [ # compare to the grid version: following lines have
                        # distributions, not values.
                        {'knn__n_neighbors': randint(2, 10),
                         'knn__p': [1, 2]
                        },
                        ]

random_search = RandomizedSearchCV(knn_pipeline, param_distributions,
n_iter=5, cv=3,
                                scoring='balanced_accuracy',
                                random_state=42)
random_search.fit(X_train, y_train)
print('The best parameters are ', random_search.best_params_)

random_search.best_estimator_
random_cv_res = pd.DataFrame(random_search.cv_results_)
random_cv_res.sort_values(by='mean_test_score', ascending=False,
inplace=True)
random_cv_res.filter(regex = '^param_|mean_test_score',
axis=1).head()
```

The best parameters are {'knn__n_neighbors': 4, 'knn__p': 2}

	param_knn__n_neighbors	param_knn__p	mean_test_score
2	4	2	0.562930
1	6	1	0.550195
3	6	1	0.550195
0	8	2	0.535366
4	8	2	0.535366

#####HalvingRandomSearchCV

```

from sklearn.experimental import enable_halving_search_cv
from sklearn.model_selection import HalvingRandomSearchCV

from scipy.stats import randint

param_distributions = [
    {'knn__n_neighbors': randint(2, 10),
     'knn__p': [1, 2]}
]
# This is where we switch to halving version ...
halving_random_search = HalvingRandomSearchCV(knn_pipeline,
param_distributions,
                                              n_candidates=5, cv=3,
                                              min_resources='exhaust',
                                              scoring='balanced_accuracy',
                                              random_state=42)
halving_random_search.fit(X_train, y_train)
print('The best parameters are ', halving_random_search.best_params_)

halving_random_search.best_estimator_
halving_random_cv_res =
pd.DataFrame(halving_random_search.cv_results_)
# In the end, we care about performances in the last iteration (using
most data)
# So, let's sort by iteration (descending), then by test score
(descending)
halving_random_cv_res.sort_values(by=['iter', 'mean_test_score'],
ascending=False, inplace=True)
# and check the top few rows
halving_random_cv_res.filter(regex = '(iter|^param_|mean_test_score|
n_resources)', axis=1).head()

```

The best parameters are {'knn__n_neighbors': 4, 'knn__p': 2}

	iter	n_resources	param_knn__n_neighbors	param_knn__p	mean_test_score
6	1	59010	4	2	0.562930
5	1	59010	6	1	0.550195

2	0	19670	4	2
0.521509				
1	0	19670	6	1
0.512701				
3	0	19670	6	1
0.512701				

####Model Retraining

```
knn_pipeline_new = Pipeline([
    ('preprocessing', pipeline),
    ('knn', KNeighborsClassifier(p=1, n_neighbors=2))
])
```

####Model Reevaluation

```
knn_pipeline_new.fit(X_train, y_train)
y_pred = knn_pipeline_new.predict(X_test)
```


#####Balanced Accuracy

```
calculate_accuracy(y_test, y_pred)
Accuracy=0.9465, Balanced Accuracy=0.6341
```

#####Confusion Matrix

```
c_matrix(y_test, y_pred, knn_pipeline_new)
confusion matrix=

array([[18253,    79],
       [  974,   365]])
```



#####Classification Report


```
class_report(y_test, y_pred)
```

```
classification report=
```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	18332
1	0.82	0.27	0.41	1339
accuracy			0.95	19671
macro avg	0.89	0.63	0.69	19671
weighted avg	0.94	0.95	0.93	19671

#####ROC Curve

```
roc_curve(knn_pipeline_new, X_test, y_test)
```

#####Cost Matrix

```
cost_matrix(knn_pipeline_new, X_train, y_train)
average cost for simple pipeline= 5977.0
```

####Techniques for Dataset Imbalance

#####K-fold

```
# k-fold
from sklearn.model_selection import cross_val_score
knn_scores = cross_val_score(knn_pipeline_new, X_train, y_train, cv=5,
scoring='balanced_accuracy')
print(f'The balanced accuracy of K Nearest Neighbor is
{knn_scores.mean():.3f}.')
```

The balanced accuracy of K Nearest Neighbor is 0.616.

#####Bootstrap

```
# Bootstrap
# Import necessary libraries
from sklearn.utils import resample
```

```

from sklearn.metrics import balanced_accuracy_score
import numpy as np

def bootstrap(model, X_train, y_train, X_test, y_test,
n_iterations=100):
    accuracy_scores = []

    for times in range(n_iterations):
        # Bootstrap resampling on the minority class
        X_resample, y_resample = resample(X_train[y_train == 1],
y_train[y_train == 1], replace=True, n_samples=X_train[y_train ==
0].shape[0])
        X_resample = np.concatenate((X_resample, X_train[y_train ==
0]), axis=0) # Combine minority and majority class along axis 0
        y_resample = np.concatenate((y_resample, y_train[y_train ==
0]), axis=0)
        X_resample = pd.DataFrame(X_resample, columns=['Current
assets',
'Cost of goods sold', 'Depreciation and amortization',
'EBITDA',
'Inventory', 'Net Income', 'Total Receivables', 'Market Value',
'Net Sales', 'Total Assets', 'Total Long-term Debt', 'EBIT',
'Gross Profit', 'Total Current Liabilities', 'Retained
Earnings',
'Total Revenue', 'Total Liabilities', 'Total Operating
Expenses'])
        y_resample = pd.DataFrame(y_resample,
columns=['is_bankruptcy'])

        model.fit(X_resample, y_resample)
        y_pred = model.predict(X_test)
        accuracy = balanced_accuracy_score(y_test, y_pred)
        accuracy_scores.append(accuracy)

    return np.array(accuracy_scores)

bootstrap_scores = bootstrap(knn_pipeline_new, X_train, y_train,
X_test, y_test)
print("Balanced accuracy after using bootstrap:",
bootstrap_scores.mean())

```

Balanced accuracy after using bootstrap: 0.7146843108041098

#####SMOTE Technique

```

#smote
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline, make_pipeline
smote = SMOTE(random_state=42)

```

```
# Create the pipeline
knn_smote_pipe = Pipeline([
    ('preprocessing', pipeline),
    ('smote', smote),
    ('knn', KNeighborsClassifier(p=1, n_neighbors=2))
])

knn_smote_pipe.fit(X_train, y_train)
y_pred2_smote = knn_smote_pipe.predict(X_test)
accuracy = accuracy_score(y_test, y_pred2_smote)
balanced_accuracy = balanced_accuracy_score(y_test, y_pred2_smote)
print(f'Accuracy={accuracy:.4f}, Balanced
Accuracy={balanced_accuracy:.4f}')

Accuracy=0.9178, Balanced Accuracy=0.7752
```

The K Nearest Neighbors (KNN) model demonstrated remarkable performance improvements. Initially, the balanced accuracy stood at 0.603, reflecting competent but not exceptional results. Through the strategic application of the SMOTE technique to address imbalanced data concerns, the balanced accuracy surged to an impressive 0.7659.

This significant increase in balanced accuracy underscores the efficacy of leveraging advanced techniques to handle class imbalances, thereby enhancing the predictive capabilities of the KNN model. The outcome reinforces the model's adaptability and its ability to discern patterns effectively, ultimately contributing to more robust and accurate predictions.

7.4 SVM

Model Training

```
from sklearn.svm import SVC

c = 1 # @param {type:"slider", min:1, max:10}

lin_svc = SVC(C=c, kernel='linear', degree=1)

svm_pipeline = Pipeline([
    ("preprocessing", pipeline),
    ("svm", lin_svc),
])

svm_pipeline.fit(X_train, y_train)
print(lin_svc.n_support_)

svm_pipeline
[4412 3881]

Pipeline(steps=[('preprocessing',
                  ColumnTransformer(transformers=[('num',
```

```

Pipeline(steps=[('imputer',
SimpleImputer(strategy='median')),
('scaler',
StandardScaler())]),
<sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ada0>),
('cat',
Pipeline(steps=[('imputer',
SimpleImputer(strategy='most_frequent')),
('cat_encoder',
OneHotEncoder())]),
<sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ace0>)]),
('svm', SVC(C=1, degree=1, kernel='linear'))])

# Assuming svm_pipeline is already defined and fitted
y_pred_svm = svm_pipeline.predict(X_test)

```

Model Evaluation

Balanced Accuracy

```

calculate_accuracy(y_test, y_pred_svm)

Accuracy=0.9319, Balanced Accuracy=0.5000

```

Confusion Matrix


```

c_matrix(y_test, y_pred_svm, svm_pipeline)

confusion matrix=

array([[18332,    0],
       [ 1339,    0]])

```



Classification Report


```
class_report(y_test, y_pred_svm)
```

```
classification report=
```

	precision	recall	f1-score	support
0	0.93	1.00	0.96	18332
1	0.00	0.00	0.00	1339
accuracy			0.93	19671
macro avg	0.47	0.50	0.48	19671
weighted avg	0.87	0.93	0.90	19671

ROC Curve

```
roc_curve(svm_pipeline, X_test, y_test)
```



Cost Matrix

```
cost_matrix(svm_pipeline, X_train, y_train)
average cost for simple pipeline= 7762.0
```

Parameter Tuning

Grid Search

```
from sklearn.model_selection import GridSearchCV

svm_pipeline = Pipeline([
    ("preprocessing", pipeline),
    ("svm", SVC()),
])

param_grid = [
    {'svm__kernel': ['linear'], 'svm__C': [1, 10, 100, 1000]},
    {'svm__kernel': ['rbf'], 'svm__C': [1, 10, 100, 1000], 'svm__gamma':
[0.001, 0.0001]},
    {'svm__kernel': ['poly'], 'svm__C': [1, 10, 100, 1000],
'svm__gamma': [0.001, 0.0001], 'svm__degree': [2, 3, 4]},
]

grid_search = GridSearchCV(svm_pipeline, param_grid, cv=3,
```

```

scoring='accuracy')
grid_search.fit(X_train, y_train)

grid_cv_res = pd.DataFrame(grid_search.cv_results_)
grid_cv_res.sort_values(by="mean_test_score", ascending=False,
inplace=True)
grid_cv_res.filter(regex = '^(^param_|mean_test_score)', axis=1)

```

HalvingGridSearchCV

```

from sklearn.experimental import enable_halving_search_cv
from sklearn.model_selection import HalvingGridSearchCV
from sklearn.svm import SVC
from sklearn.pipeline import make_pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
import numpy as np
import pandas as pd

param_grid = [{
    'svm__kernel': ['linear', 'poly', 'rbf', 'sigmoid'],
    'svm__C': [0.1, 1, 10],
    'svm__gamma': ['scale', 'auto', 0.1, 1]
}]

# Check what's in this parameter grid
print('The parameter grid: ')
print(param_grid)

# Change to the new strategy starting here
halving_grid_search = HalvingGridSearchCV(svm_pipeline, param_grid,
cv=3,
                                         min_resources='exhaust', #
                                         use_all_data_in_the_last_round, start_with_as_needed
                                         scoring='balanced_accuracy')
halving_grid_search.fit(X_train, y_train)

The parameter grid:
[{'svm__kernel': ['linear', 'poly', 'rbf', 'sigmoid'], 'svm__C': [0.1,
1, 10], 'svm__gamma': ['scale', 'auto', 0.1, 1]}]

HalvingGridSearchCV(cv=3,
                    estimator=Pipeline(steps=[('preprocessing',
ColumnTransformer(transformers=[('num',
Pipeline(steps=[('imputer',
SimpleImputer(strategy='median')),
('scaler',

```

```

StandardScaler()))],

<sklearn.compose._column_transformer.make_column_selector object at
0x7dd9e061b6a0>),

('cat',

Pipeline(steps=[('imputer',

SimpleImputer(strategy='most_frequent')),

('cat_encoder',

OneHotEncoder()))],

<sklearn.compose._column_transformer.make_column_selector object at
0x7dd9e061b580>)])),

('svm',
 SVC(C=1, degree=1,

kernel='linear'))],

param_grid=[{'svm__C': [0.1, 1, 10],
              'svm__gamma': ['scale', 'auto', 0.1,
1],
              'svm__kernel': ['linear', 'poly',
'rbf',
              'sigmoid']}],

scoring='balanced_accuracy')

print('The best parameters are ', halving_grid_search.best_params_)

halving_grid_cv_res = pd.DataFrame(halving_grid_search.cv_results_) #
convert to DF for convenience
# In the end, we care about performances in the last iteration (using
most data)
# So, let's sort by iteration (descending), then by test score
(descending)
halving_grid_cv_res.sort_values(by=['iter', 'mean_test_score'],
ascending=False, inplace=True)
# and check the top few rows
halving_grid_cv_res.filter(regex='(iter|^param_|mean_test_score|
n_resources)', axis=1).head()

The best parameters are {'svm__C': 1, 'svm__gamma': 'auto',
'svm__kernel': 'rbf'}

  iter  n_resources  param_svm__C  param_svm__gamma  param_svm__kernel
\
71    3         58995             1             auto             rbf

```


70	3	58995	0.1	scale	linear
65	2	19665	1	auto	rbf
64	2	19665	1	auto	poly
66	2	19665	1	0.1	linear

	mean_test_score
71	0.502532
70	0.500000
65	0.501478
64	0.500000
66	0.500000

RandomSearchCV

```

from sklearn.svm import SVC
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import reciprocal, uniform

# Define the parameter distributions for RandomizedSearchCV
param_grid = [{
    'svm__kernel': ['linear', 'poly', 'rbf', 'sigmoid'],
    'svm__C': [0.1, 1, 10],
    'svm__gamma': ['scale', 'auto', 0.1, 1]
}]

# Perform RandomizedSearchCV
random_search = RandomizedSearchCV(svm_pipeline,
param_distributions=param_grid, n_iter=5, cv=3,
                                scoring='balanced_accuracy',
random_state=42)
random_search.fit(X_train, y_train)

# Print the best parameters
print('The best parameters are ', random_search.best_params_)

# Get the best estimator
best_svm_estimator = random_search.best_estimator_

# Access other details if needed
random_cv_res = pd.DataFrame(random_search.cv_results_)
random_cv_res.sort_values(by='mean_test_score', ascending=False,
inplace=True)
random_cv_res.filter(regex='(^param_|mean_test_score)', axis=1).head()

```

HalvingRandomSearchCV

```
from sklearn.model_selection import HalvingRandomSearchCV
from sklearn.svm import SVC
from scipy.stats import reciprocal, uniform

# Define the parameter distributions for HalvingRandomSearchCV
param_distributions = {
    'svm__kernel': ['linear', 'poly', 'rbf', 'sigmoid'],
    'svm__C': reciprocal(0.1, 10),
    'svm__gamma': reciprocal(0.01, 1.0)
}

# Perform HalvingRandomSearchCV
halving_random_search = HalvingRandomSearchCV(svm_pipeline,
param_distributions=param_distributions,
                                                factor=3,
n_resources=100, random_state=42)
halving_random_search.fit(X_train, y_train)

# Print the best parameters
print('The best parameters are ', halving_random_search.best_params_)

# Get the best estimator
best_svm_estimator = halving_random_search.best_estimator_

# Access other details if needed
halving_random_cv_res =
pd.DataFrame(halving_random_search.cv_results_)
halving_random_cv_res.sort_values(by='mean_test_score',
ascending=False, inplace=True)
halving_random_cv_res.filter(regex='(^param_|mean_test_score)',
axis=1).head()
```

Model Retraining

```
from sklearn.svm import SVC

lin_svc = SVC(C=1, kernel='rbf', gamma='auto')

svm_pipeline = Pipeline([
    ("preprocessing", pipeline),
    ("svm", lin_svc),
])

svm_pipeline.fit(X_train, y_train)
#print(lin_svc.n_support_)
svm_pipeline

Pipeline(steps=[('preprocessing',
                  ColumnTransformer(transformers=[('num',
```

```

Pipeline(steps=[('imputer',
SimpleImputer(strategy='median')),
('scaler',
StandardScaler())]),
<sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ada0>),
('cat',
Pipeline(steps=[('imputer',
SimpleImputer(strategy='most_frequent')),
('cat_encoder',
OneHotEncoder())]),
<sklearn.compose._column_transformer.make_column_selector object at
0x77ff6f07ace0>)]),
('svm', SVC(C=1, gamma='auto'))])

# Assuming svm_pipeline is already defined and fitted
y_pred_svm_new = svm_pipeline.predict(X_test)

```

Model Reevaluation

Balanced Accuracy

```

calculate_accuracy(y_test, y_pred_svm_new)

Accuracy=0.9323, Balanced Accuracy=0.5030

```

Confusion Matrix


```

c_matrix(y_test, y_pred_svm_new, svm_pipeline)

confusion matrix=

array([[18332,    0],
       [ 1331,    8]])

```



Classification Report


```
class_report(y_test, y_pred_svm_new)
```

```
classification report=
```

	precision	recall	f1-score	support
0	0.93	1.00	0.96	18332
1	1.00	0.01	0.01	1339
accuracy			0.93	19671
macro avg	0.97	0.50	0.49	19671
weighted avg	0.94	0.93	0.90	19671

ROC Curve

```
roc_curve(svm_pipeline, X_test, y_test)
```



Cost Matrix

```
cost_matrix(svm_pipeline, X_train, y_train)
average cost for simple pipeline= 7711.4
```

Techniques for Dataset Imbalance

#####K-fold

```
from sklearn.model_selection import cross_val_score
svm_scores = cross_val_score(svm_pipeline, X_train, y_train, cv=5,
                              scoring='balanced_accuracy')
print(f'The balanced accuracy of K Nearest Neighbor is
{svm_scores.mean():.3f}.')
```

The balanced accuracy of K Nearest Neighbor is 0.503.

Balanced Weights

```
from sklearn.svm import SVC
from sklearn.pipeline import make_pipeline

# Modify the pipeline to include balanced class weights
modified_svm_pipeline = svm_pipeline
modified_svm_pipeline.named_steps['svm'].set_params(class_weight='bala
```

```

nced')

# Fit the modified pipeline on training data
modified_svm_pipeline.fit(X_train, y_train)

# Make predictions on the test set
y_pred = modified_svm_pipeline.predict(X_test)

calculate_accuracy(y_test, y_pred)

Accuracy=0.2905, Balanced Accuracy=0.5864

```

SMOTE Technique

```

from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline, make_pipeline

svc_smote_pipe = make_pipeline(pipeline, SMOTE(random_state=42),
SVC(C=1, kernel='rbf', gamma='auto'))
svc_smote_pipe
svc_smote_pipe.fit(X_train, y_train)
y_pred = svc_smote_pipe.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
balanced_accuracy = balanced_accuracy_score(y_test, y_pred)
print(f'Accuracy={accuracy:.4f}, Balanced
Accuracy={balanced_accuracy:.4f}')

Accuracy=0.3050, Balanced Accuracy=0.5904

```

The performance of the SVM model in predicting our dataset proved to be less than satisfactory. Despite exhaustive efforts in hyperparameter tuning, the balanced accuracy remained stagnant at 0.50, with the best parameters identified as $C=1$, $\text{kernel}='rbf'$, and $\text{gamma}='auto'$ through the Halving Grid Search Model.

While alternative search parameters were considered for SVM, the iterative exploration encountered challenges, notably infinite loops and considerable time consumption. As a pragmatic solution, we opted for the Halving Grid Search, prioritizing efficiency and practicality.

Despite diligent efforts in addressing class imbalance, the SVM model exhibited limited improvement in balanced accuracy. This observation underscores the inherent complexities and challenges associated with optimizing SVM for our specific dataset. Nevertheless, our commitment to exploring various avenues for improvement remains steadfast, reflecting the dynamic and iterative nature of model refinement.

###7.5 Voting

####Balanced Accuracy

```


from sklearn.ensemble import VotingClassifier

```

```
clf_voting = VotingClassifier(  
    estimators=[  
        ('lr', logis_smote_pipe),  
        ('rf', rf_smote_pipe),  
        ('knn', knn_smote_pipe),  
        ('svm', svc_smote_pipe) ])  
  
clf_voting.fit(X_train, y_train)  
y_pred_voting = clf_voting.predict(X_test)  
calculate_accuracy(y_test, y_pred_voting)  
  
Accuracy=0.8910, Balanced Accuracy=0.7643
```

####Confusion Matrix

```
c_matrix(y_test, y_pred_voting, clf_voting)  
confusion matrix=  
  
array([[16700,  1632],  
       [  512,   827]])
```



####Classification Report

```
class_report(y_test, y_pred_voting)
```

	precision	recall	f1-score	support
0	0.97	0.91	0.94	18332
1	0.34	0.62	0.44	1339
accuracy			0.89	19671
macro avg	0.65	0.76	0.69	19671
weighted avg	0.93	0.89	0.91	19671

####Cost Matrix


```
cost_matrix(clf_voting, X_train, y_train)
```

average cost for simple pipeline= 4172.0


###Final Chosen Model Reevaluation(KNN-SMOTE)

To select the final model according to our results for better understanding we have plotted graphs representing the Balanced accuracy and accuracy of various models.

The graphical representation below offers insights into the accuracy and balanced accuracies of models prior to undergoing Hyperparameter Tuning and efforts to address class imbalances. The conspicuous gap observed in the accuracies underscores the pronounced impact of class imbalance on model performance.



Moving ahead, the selection of the optimal model for predictions became paramount. The graph below illuminates the balanced accuracy of the retrained model, comparing its performance before and after the application of class imbalance techniques.



Analyzing the graph it is evident that the KNN SMOTE is the best Model so far to move ahead with for the predictions and our analysis as it gives us the best Balanced Accuracy.


####Balanced Accuracy

```
calculate_accuracy(y_test, y_pred2_smote)
```

Accuracy=0.9178, Balanced Accuracy=0.7752

####Confusion Matrix

```
c_matrix(y_test, y_pred2_smote, knn_smote_pipe)
confusion matrix=
array([[17237, 1095],
       [ 522,  817]])
```



####Classification Report


```
class_report(y_test, y_pred2_smote)
classification report=
```

	precision	recall	f1-score	support
0	0.97	0.94	0.96	18332
1	0.43	0.61	0.50	1339
accuracy			0.92	19671

macro avg	0.70	0.78	0.73	19671
weighted avg	0.93	0.92	0.92	19671

#####ROC Curve

```
roc_curve(knn_smote_pipe, X_test, y_test)
```



#####Cost Matrix

```
cost_matrix(knn_smote_pipe, X_train, y_train)
average cost for simple pipeline= 3960.8
```

###7.6 Final Prediction

```
y_pred2_smote2 = knn_smote_pipe.predict(X)
df1 = df.copy()
df1['predict_bankrupt'] = y_pred2_smote2
df1 # df1 = add the new prediction on the original dataset(df)
```

	is_bankruptcy	Current assets	Cost of goods sold
\			
company_name	year		

C_1	1999	0	511.267	833.107
	2000	0	485.856	713.811
	2001	0	436.656	526.477
	2002	0	396.412	496.747
	2003	0	432.204	523.302
...
C_8971	2014	0	233.211	43.338
	2015	0	105.559	59.184
	2016	0	63.971	69.074
	2017	0	135.207	66.527
	2018	0	82.589	68.817
Inventory \ Depreciation and amortization EBITDA				
company_name	year			
C_1	1999	18.373	89.031	336.018
	2000	18.577	64.367	320.590
	2001	22.496	27.207	286.588
	2002	27.172	30.745	259.954
	2003	26.680	47.491	247.245
...
C_8971	2014	14.094	45.615	3.376
	2015	42.592	202.133	2.288
	2016	65.057	79.051	2.581
	2017	65.330	69.171	2.013
	2018	65.201	67.262	2.112
Net Income Total Receivables Market Value Net				

Sales \				
company_name	year			
C_1	1999	35.163	128.348	372.7519
1024.333	2000	18.531	115.187	377.1180
874.255	2001	-58.939	77.528	364.5928
638.721	2002	-12.410	66.322	143.3295
606.337	2003	3.504	104.661	308.9071
651.958
...
C_8971	2014	25.261	22.846	756.4827
104.223	2015	129.688	54.611	527.5750
291.153	2016	-1.442	42.467	578.8868
169.858	2017	-20.401	27.217	412.6241
161.884	2018	-50.946	45.839	354.1786
160.513				
		Total Assets	Total Long-term Debt	EBIT Gross
Profit \				
company_name	year			
C_1	1999	740.998	180.447	70.658
191.226	2000	701.854	179.987	45.790
160.444	2001	710.199	217.699	4.711
112.244	2002	686.621	164.658	3.573
109.590	2003	709.292	248.666	20.811
128.656
...
C_8971	2014	1099.101	184.666	31.521
60.885	2015	1865.926	770.103	159.541
231.969	2016	1746.235	683.985	13.994
100.784	2017	1736.110	694.035	3.841
95.357				

91.696	2018	1625.370	632.122	2.061
Total Current Liabilities Retained Earnings \				
company_name	year			
C_1	1999	163.816	201.026	
	2000	125.392	204.065	
	2001	150.464	139.603	
	2002	203.575	124.106	
	2003	131.261	131.884	
...		...		
C_8971	2014	28.197	28.095	
	2015	88.128	157.783	
	2016	85.765	156.341	
	2017	82.010	135.941	
	2018	79.365	84.995	
Total Revenue Total Liabilities Total Operating				
Expenses \	company_name	year		
C_1	1999	1024.333	401.483	
935.302				
	2000	874.255	361.642	
809.888				
	2001	638.721	399.964	
611.514				
	2002	606.337	391.633	
575.592				
	2003	651.958	407.608	
604.467				
...		
...				
C_8971	2014	104.223	225.887	
58.608				
	2015	291.153	880.327	
89.020				
	2016	169.858	770.233	
90.807				
	2017	161.884	776.697	
92.713				
	2018	160.513	712.687	
93.251				
predict_bankrupt				
company_name	year			
C_1	1999	0		
	2000	0		
	2001	0		
	2002	0		

	2003	0
...		...
C_8971	2014	0
	2015	0
	2016	0
	2017	0
	2018	0

[78682 rows x 20 columns]

```

alivedf1 = df1[df1['predict_bankrupt']==0]
faileddf1 = df1[df1['predict_bankrupt']==1]
print('basic statistics about the alive company dataset (include =
"number"):')
display(alivedf1.describe(include = 'number'))
print('basic statistics about the failed company dataset (include =
"number"):')
display(faileddf1.describe(include = 'number'))

```

basic statistics about the alive company dataset (include = "number"):

	is_bankruptcy	Current assets	Cost of goods sold \
count	72890.000000	72890.000000	72890.000000
mean	0.007175	919.641944	1655.249681
std	0.084403	4066.882328	9242.547560
min	0.000000	-7.760000	-366.645000
25%	0.000000	19.289250	17.151500
50%	0.000000	103.392500	104.322000
75%	0.000000	453.069000	658.140000
max	1.000000	169662.000000	374623.000000

	Depreciation and amortization	EBITDA	Inventory \
count	72890.000000	72890.000000	72890.000000
mean	124.386420	396.495107	209.889237
std	673.916692	2084.059773	1094.011669
min	0.000000	-21913.000000	0.000000
25%	1.188000	-0.610000	0.000000
50%	7.971000	15.822000	7.205000
75%	48.235750	145.798750	77.602500
max	28430.000000	81730.000000	62567.000000

	Net Income	Total Receivables	Market Value	Net
Sales \				
count	72890.000000	72890.000000	7.289000e+04	72890.000000
mean	143.498819	299.481590	3.618874e+03	2461.642678
std	1303.855903	1382.600192	1.909718e+04	12370.953151
min	-98696.000000	-0.006000	1.000000e-04	-1964.999000

25%	-6.287000	3.420000	3.594120e+01	28.018750
50%	2.157000	23.626500	2.426296e+02	188.548500
75%	44.062750	137.067750	1.334322e+03	1092.605750
max	104821.000000	65812.000000	1.073391e+06	511729.000000

	Total Assets	Total Long-term Debt	EBIT	Gross Profit \
count	72890.000000	72890.000000	72890.000000	72890.000000
mean	2967.515471	733.204506	272.108542	806.393111
std	13337.404277	3318.835009	1547.672372	3910.838492
min	0.001000	-0.023000	-25913.000000	-21536.000000
25%	37.859500	0.000000	-2.333750	8.915000
50%	216.136000	7.057500	7.360500	65.436000
75%	1207.959250	246.048000	93.300000	359.248000
max	531864.000000	166250.000000	71230.000000	137106.000000

	Total Current Liabilities	Retained Earnings	Total Revenue \
count	72890.000000	72890.000000	72890.000000
mean	632.273175	586.744846	2461.642678
std	3032.882600	6596.012562	12370.953151
min	0.001000	-102362.000000	-1964.999000
25%	8.827250	-62.774750	28.018750
50%	43.572000	-0.002500	188.548500
75%	230.761250	161.432250	1092.605750
max	116866.000000	402089.000000	511729.000000

	Total Liabilities	Total Operating Expenses	predict_bankrupt
count	72890.000000	72890.000000	72890.0
mean	1818.429628	2065.147598	0.0
std	8286.054262	10782.557289	0.0
min	0.001000	-317.197000	0.0
25%	13.351000	33.075000	0.0
50%	81.200000	171.355500	0.0
75%	637.286000	911.418500	0.0
max	337980.000000	481580.000000	0.0

basic statistics about the failed company dataset (include = "number"):

	is_bankruptcy	Current assets	Cost of goods sold \		
count	5792.000000	5792.000000	5792.000000		
mean	0.810946	386.046215	830.383940		
std	0.391585	1119.295540	2785.079876		
min	0.000000	0.001000	-0.666000		
25%	1.000000	14.995500	15.802000		
50%	1.000000	74.146000	96.707500		
75%	1.000000	258.937000	444.781000		
max	1.000000	16103.000000	40683.000000		
	Depreciation and amortization		EBITDA	Inventory	Net
Income \					
count	5792.000000	5792.000000	5792.000000		
5792.000000					
mean	81.565534	128.393766	97.360935	-	
48.266364					
std	253.758486	518.106680	460.458992		
573.049955					
min	0.000000	-8218.500000	0.000000	-	
27446.000000					
25%	1.239000	-5.013500	0.029250	-	
27.060000					
50%	7.590000	5.929000	5.456000	-	
3.882000					
75%	45.370250	79.491750	48.057000		
8.092250					
max	5475.000000	6136.000000	9963.000000		
4504.000000					
	Total Receivables	Market Value	Net Sales	Total Assets \	
count	5792.000000	5792.000000	5792.000000	5792.000000	
mean	127.652074	840.555791	1135.473712	1603.555791	
std	403.122021	3078.639202	3515.478844	5153.768371	
min	0.000000	0.001300	-0.234000	0.005000	
25%	2.139250	26.403475	21.116250	32.222000	
50%	14.266000	118.443050	160.614000	183.981000	
75%	78.459750	504.684350	678.461500	811.636000	
max	6786.000000	139092.655000	53012.000000	76995.000000	
	Total Long-term Debt	EBIT	Gross Profit \		
count	5792.000000	5792.000000	5792.000000		
mean	587.566787	46.828232	305.089772		
std	2040.308534	395.878152	921.392672		
min	0.000000	-10537.000000	-8001.000000		
25%	0.175750	-10.437250	4.256000		
50%	16.480500	0.030000	42.845500		
75%	278.401000	35.271750	217.737750		
max	21586.000000	4822.000000	15192.000000		
	Total Current Liabilities	Retained Earnings	Total Revenue \		

count	5792.000000	5792.000000	5792.000000
mean	330.682565	-150.597021	1135.473712
std	1203.856332	1746.390517	3515.478844
min	0.005000	-43091.000000	-0.234000
25%	9.554750	-139.146750	21.116250
50%	40.277500	-27.371500	160.614000
75%	154.569250	33.584750	678.461500
max	41695.000000	7832.000000	53012.000000

	Total Liabilities	Total Operating Expenses	predict_bankrupt
count	5792.000000	5792.000000	5792.0
mean	1208.947542	1007.079946	1.0
std	4091.377077	3270.758436	0.0
min	0.005000	-0.016000	1.0
25%	15.373750	30.450000	1.0
50%	90.822500	143.558500	1.0
75%	548.848000	580.510250	1.0
max	64092.000000	49363.000000	1.0

```

a = alivedf1.describe(include = 'number').iloc[1] # a dataset
represent the alive company info only display mean result
b = faileddf1.describe(include = 'number').iloc[1] # b dataset
represent the alive company info only display mean result
a.drop(['is_bankruptcy', 'predict_bankrupt'], inplace=True)
b.drop(['is_bankruptcy', 'predict_bankrupt'], inplace=True)

import matplotlib.pyplot as plt
import numpy as np


features = a.index
values_a = a.values
values_b = b.values

bar_width=0.35
index = np.arange(len(features))

plt.figure(figsize=(12, 8))
plt.bar(index, values_a , bar_width, label='alive',color = 'blue')
plt.bar(index + bar_width, values_b , bar_width, label='failed',color
= 'red')
plt.xlabel('Features')
plt.ylabel('Values')
plt.title('Comparison of Features between different bankrupt
condition')
plt.xticks(index + bar_width / 2, features, rotation='vertical')
plt.legend()
plt.tight_layout()

plt.show()

```



Insights

From financial metrics, the mean current assets for alive companies stand substantially higher than for failed companies, pointing to successful companies maintaining elevated levels of liquid assets, a contributing factor to their financial stability. In contrast, alive companies exhibit positive net income on average, whereas failed companies show a negative net income, underscoring the financial losses experienced by failed companies leading to bankruptcy.

Surprisingly, our anticipation regarding higher long-term debt and total liabilities for failed companies is contradicted. This suggests that successful companies may be more open to leveraging long-term financing for strategic growth or investments. Moreover, these successful companies adeptly manage financial obligations to sustain their operations. Operationally, alive companies boast higher EBITDA, indicative of robust cash flows from operational activities. However, they also incur higher operating expenses, hinting at the costs associated with maintaining successful operations.

From a market value perspective, alive companies exhibit substantially higher values compared to failed companies, signaling heightened investor confidence in the future potential and performance of successful companies. These insights drive our motivation throughout the project and hold importance for stakeholders, investors, and financial institutions seeking a nuanced understanding of the risks associated with companies' financial health and bankruptcy potential.

##8. Challenge

Given the highly imbalanced nature of our dataset, where approximately 97% of datapoints represent the "alive" class and only 3% represent the "failed" class, we implement several strategies for effective evaluation.

Firstly, in response to the challenge posed by imbalanced datasets, we prioritize balanced accuracy over traditional accuracy metrics when assessing our models. Traditional accuracy can be misleading in imbalanced scenarios as it tends to favor the majority class. Balanced accuracy, which considers the mean of sensitivity and specificity, offers a more reliable evaluation across all classes.

Secondly, in our model tuning, we incorporate the `class_weight='balanced'` hyperparameter to address the impact of imbalanced data. This adjustment ensures that the model assigns appropriate weights to different classes during training, giving emphasis to the minority class and enhancing predictive accuracy for both major and minor classes.

Thirdly, we employ resampling techniques, specifically over-sampling, to tackle the imbalance. This involves duplicating or generating synthetic samples for the minority class, providing the model with more instances to learn and distinguish patterns related to that class. This technique is crucial in preventing bias towards the majority class and fostering a more equitable model.

Fourthly, we also leverage the SMOTE (Synthetic Minority Over-sampling Technique) method to address the imbalanced problem. SMOTE helps by creating new data points for the smaller group, making the data more balanced. This way, the model can learn better from both groups and make fair predictions.

Lastly, K-fold cross-validation serves as an additional tool to address this problem. Specifically, K-fold cross-validation, when combined with these strategies, helps us navigate the challenges posed by imbalanced datasets.

Through these thoughtful approaches, our goal is to enhance the robustness and fairness of our machine learning models in the context of imbalanced datasets.

##9. Conclusion

In our comprehensive project journey, our primary objective was to craft a classification model that excels in balanced accuracy, particularly within the intricate landscape of a highly imbalanced dataset. Our exploration spanned various models, encompassing Logistic Regression, Random Forest, K-Nearest Neighbors (KNN), and Support Vector Machines (SVM). Despite exhaustive efforts in training diverse models, the incremental gains in balanced accuracy were modest.

Beyond model selection, our focus extended to feature selection, delving into the nuanced impact of each variable on shaping the growth and success trajectories of the companies under scrutiny.

In summary, our project underscores the paramount importance of achieving peak accuracy in predictive modeling. The preference for KNN SMOTE signifies its adeptness in navigating imbalanced data challenges. Additionally, insights gained from feature selection illuminate how specific variables shape the growth dynamics of companies.

Further enriching our understanding is the recognition that stakeholders hold different perspectives on particular issues. For companies, a grasp of their financial health guides future

plans, influencing decisions to earn more, spend judiciously, or seek external assistance. Simultaneously, investors keenly analyze a company's financial standing to make informed investment decisions, directing capital where growth potential is perceived. Predicting a company's trajectory, therefore, becomes pivotal for both business entities and investors alike, shaping strategic actions and investment choices. Ultimately, our work not only aids investors in making well-informed decisions but also empowers companies to navigate toward success with a nuanced understanding of their growth dynamics.